

Cifar_10_Final_Project

December 13, 2019

1 Final Project - Cifar 10

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. We will be building a model that predicts/classifies what classes the images fall into. For this project, we will be using Keras.

1.1 Getting Data

To begin, we will load the cifar-10 dataset. `X_train` and `y_train` will contain information for 50,000 training images. `X_test` and `y_test` will contain information for 10,000 test images. We will leave the test images alone and only test on the model that produces the best result on the validation set. We also create an array that labels what the 10 classes should be.

Additionally, we will create a validation set once we train the model, randomly assigning 20% of the data to be validation data.

```
[0]: from tensorflow import keras
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from keras.utils import print_summary, to_categorical
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.regularizers import l2
from keras import regularizers

import numpy as np
import os

import pandas as pd
from sklearn.metrics import classification_report, confusion_matrix

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

label_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

1.2 Useful Helper Functions

Below, we will create a few helper functions that will help us evaluate our model.

1.2.1 Validation Helper

This helper function will help plot the training and validation loss graph to see if we're underfitting or overfitting the data. We will also plot the training and validation accuracy data to compare with the loss.

1.2.2 Test Set Evaluation

The first function (`get_class_from_softmax`) will help us turn an array of probabilities of likelihood of each class (`softmax_list`) to a the index of the maximum probability, which we can convert to a class using the `label_names` array (`convert_to_labels` function).

We will then use the helpers above in the `get_metrics` function, which will produce a classification report and confusion matrix for the test set.

```
[0]: # Training vs validation loss graph

import matplotlib.pyplot as plt

def plot_graphs(history):
    # Plot training & validation loss values
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')
    plt.show()

    # Plot training & validation accuracy values
    plt.plot(history.history['acc'])
    plt.plot(history.history['val_acc'])
    plt.title('Model accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')
    plt.show()
```

```
[0]: ## Here we'll use the model to classify the test set!

import pandas as pd

def get_class_from_softmax(softmax_list):
    maximum = 0
    max_index = 0
```

```

    for i in range(len(softmax_list)):
        if softmax_list[i] > maximum:
            maximum = softmax_list[i]
            max_index = i

    return max_index

def convert_to_labels(A):
    labels = list()
    for row in A:
        labels.append(label_names[get_class_from_softmax(row)])
    return labels

def get_metrics(model_name):
    model = load_model(model_name)

    print(model.evaluate(x=x_test, y=to_categorical(y_test, num_classes)))

    predicted = model.predict(x_test)

    predicted_labels = convert_to_labels(predicted)
    actual_labels = convert_to_labels(to_categorical(y_test, num_classes))

    # Generate classification report
    print(classification_report(y_true=actual_labels,
                               y_pred=predicted_labels,
                               labels=label_names))

    # Generate the confusion matrix
    matrix = pd.DataFrame(
        confusion_matrix(actual_labels, predicted_labels, labels=label_names),
        index=label_names,
        columns=['predicted - airplane',
                 'predicted - automobile',
                 'predicted - bird',
                 'predicted - cat',
                 'predicted - deer',
                 'predicted - dog',
                 'predicted - frog',
                 'predicted - horse',
                 'predicted - ship',
                 'predicted - truck']
    )
    pd.set_option('display.max_columns', None)

    print(matrix)

```

1.3 Model Creating, Training, and Validating

We will begin with a batch size of 32 and 25 epochs.

Additionally, note that when we're training(fitting) the model, we add the validation split of 0.2, which splits the training data randomly into 80% training and 20% validation so we can see how well the model is performing. We will be using this split for all models below.

```
history = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)
```

```
[0]: batch_size = 32
      num_classes = 10
      epochs = 25

      shape = (32, 32, 3)
```

1.4 Model #1

For our first model, we will begin by creating a 6 layer CNN. For the optimizer, we will arbitrarily choose "RMSprop". We will test different optimizers as we proceed as well as additional layers.

```
[7]: model = Sequential()
      model.add(Conv2D(32, (3, 3), padding='same',
                      kernel_initializer='random_uniform',
                      input_shape=shape))
      model.add(Activation('relu'))
      model.add(Conv2D(32, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(MaxPooling2D(pool_size=(2, 2)))

      model.add(Conv2D(64, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(Conv2D(64, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(MaxPooling2D(pool_size=(2, 2)))

      model.add(Conv2D(128, (3, 3),
```

```

        padding='same',
        kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
        padding='same',
        kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
        optimizer=opt,
        metrics=['accuracy'])

print(model.summary())

#Train model
history = model.fit(x_train,
        to_categorical(y_train, num_classes),
        epochs=epochs,
        verbose=2,
        validation_split=0.2,
        shuffle=True)

#Save model
model.save('model_1.h5', overwrite=True)

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/keras/initializers.py:119: calling RandomUniform.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.
Instructions for updating:

If using Keras pass *_constraint arguments to layers.
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
activation_1 (Activation)	(None, 32, 32, 32)	0
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
activation_2 (Activation)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
activation_3 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
activation_4 (Activation)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
activation_5 (Activation)	(None, 8, 8, 128)	0
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 256)	524544
activation_6 (Activation)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2570
activation_7 (Activation)	(None, 10)	0
Total params: 814,122		
Trainable params: 814,122		
Non-trainable params: 0		



```
-----  
None  
Train on 40000 samples, validate on 10000 samples  
Epoch 1/25  
40000/40000 - 11s - loss: 1.5936 - acc: 0.4250 - val_loss: 1.3381 - val_acc:  
0.5228  
Epoch 2/25  
40000/40000 - 7s - loss: 1.2037 - acc: 0.5775 - val_loss: 1.1519 - val_acc:  
0.6031  
Epoch 3/25  
40000/40000 - 7s - loss: 0.9967 - acc: 0.6511 - val_loss: 0.9737 - val_acc:  
0.6571  
Epoch 4/25  
40000/40000 - 7s - loss: 0.8404 - acc: 0.7083 - val_loss: 0.9273 - val_acc:  
0.6816  
Epoch 5/25  
40000/40000 - 7s - loss: 0.7237 - acc: 0.7505 - val_loss: 0.8428 - val_acc:  
0.7139  
Epoch 6/25  
40000/40000 - 7s - loss: 0.6264 - acc: 0.7854 - val_loss: 0.7811 - val_acc:  
0.7359  
Epoch 7/25  
40000/40000 - 7s - loss: 0.5395 - acc: 0.8138 - val_loss: 0.8480 - val_acc:  
0.7216  
Epoch 8/25  
40000/40000 - 7s - loss: 0.4589 - acc: 0.8409 - val_loss: 0.7943 - val_acc:  
0.7477  
Epoch 9/25  
40000/40000 - 7s - loss: 0.3891 - acc: 0.8666 - val_loss: 0.8347 - val_acc:  
0.7446  
Epoch 10/25  
40000/40000 - 7s - loss: 0.3249 - acc: 0.8871 - val_loss: 0.8575 - val_acc:  
0.7444  
Epoch 11/25  
40000/40000 - 7s - loss: 0.2630 - acc: 0.9092 - val_loss: 0.8884 - val_acc:  
0.7497  
Epoch 12/25  
40000/40000 - 7s - loss: 0.2169 - acc: 0.9259 - val_loss: 0.9716 - val_acc:  
0.7481  
Epoch 13/25  
40000/40000 - 7s - loss: 0.1766 - acc: 0.9393 - val_loss: 1.0522 - val_acc:  
0.7475  
Epoch 14/25  
40000/40000 - 7s - loss: 0.1458 - acc: 0.9492 - val_loss: 1.0967 - val_acc:  
0.7553  
Epoch 15/25  
40000/40000 - 7s - loss: 0.1221 - acc: 0.9569 - val_loss: 1.2455 - val_acc:  
0.7461
```

```
Epoch 16/25
40000/40000 - 7s - loss: 0.1035 - acc: 0.9645 - val_loss: 1.2370 - val_acc:
0.7523
Epoch 17/25
40000/40000 - 7s - loss: 0.0901 - acc: 0.9689 - val_loss: 1.4726 - val_acc:
0.7436
Epoch 18/25
40000/40000 - 7s - loss: 0.0814 - acc: 0.9721 - val_loss: 1.4091 - val_acc:
0.7451
Epoch 19/25
40000/40000 - 7s - loss: 0.0739 - acc: 0.9738 - val_loss: 1.4897 - val_acc:
0.7474
Epoch 20/25
40000/40000 - 7s - loss: 0.0698 - acc: 0.9760 - val_loss: 1.5266 - val_acc:
0.7546
Epoch 21/25
40000/40000 - 7s - loss: 0.0641 - acc: 0.9789 - val_loss: 1.7113 - val_acc:
0.7436
Epoch 22/25
40000/40000 - 7s - loss: 0.0631 - acc: 0.9782 - val_loss: 1.6118 - val_acc:
0.7422
Epoch 23/25
40000/40000 - 7s - loss: 0.0599 - acc: 0.9791 - val_loss: 1.5860 - val_acc:
0.7538
Epoch 24/25
40000/40000 - 7s - loss: 0.0548 - acc: 0.9809 - val_loss: 1.7428 - val_acc:
0.7514
Epoch 25/25
40000/40000 - 7s - loss: 0.0535 - acc: 0.9822 - val_loss: 1.8211 - val_acc:
0.7472
```

1.4.1 Validating Model #1

We can see from the plot below that we are heavily overfitting the data. To prevent this, we will try adding regularization.

```
[10]: plot_graphs(history)
```

1.5 Model #2 - Regularization with Weight Decay

1.5.1 We will try adding a weight decay of 0.0005 (default)

```
[0]: from keras import regularizers

model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape,
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dense(num_classes))
```

```

model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_2_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 2.8412 - acc: 0.4051 - val_loss: 2.2614 - val_acc: 0.5002

Epoch 2/25

40000/40000 - 12s - loss: 2.0279 - acc: 0.5384 - val_loss: 1.8188 - val_acc: 0.5744

Epoch 3/25

40000/40000 - 13s - loss: 1.6899 - acc: 0.6036 - val_loss: 1.5477 - val_acc: 0.6357

Epoch 4/25

40000/40000 - 12s - loss: 1.4689 - acc: 0.6517 - val_loss: 1.4261 - val_acc: 0.6608

Epoch 5/25

40000/40000 - 13s - loss: 1.3099 - acc: 0.6872 - val_loss: 1.3336 - val_acc: 0.6715

Epoch 6/25

40000/40000 - 13s - loss: 1.1912 - acc: 0.7153 - val_loss: 1.2332 - val_acc: 0.7037

Epoch 7/25

40000/40000 - 13s - loss: 1.0914 - acc: 0.7416 - val_loss: 1.1738 - val_acc: 0.7123

Epoch 8/25

40000/40000 - 13s - loss: 1.0148 - acc: 0.7612 - val_loss: 1.1587 - val_acc: 0.7119

Epoch 9/25

40000/40000 - 13s - loss: 0.9502 - acc: 0.7790 - val_loss: 1.1558 - val_acc: 0.7120

Epoch 10/25
40000/40000 - 13s - loss: 0.8894 - acc: 0.7993 - val_loss: 1.0117 - val_acc: 0.7581
Epoch 11/25
40000/40000 - 12s - loss: 0.8411 - acc: 0.8103 - val_loss: 1.0348 - val_acc: 0.7503
Epoch 12/25
40000/40000 - 12s - loss: 0.7925 - acc: 0.8259 - val_loss: 1.0098 - val_acc: 0.7585
Epoch 13/25
40000/40000 - 13s - loss: 0.7529 - acc: 0.8389 - val_loss: 0.9925 - val_acc: 0.7647
Epoch 14/25
40000/40000 - 13s - loss: 0.7141 - acc: 0.8512 - val_loss: 0.9833 - val_acc: 0.7700
Epoch 15/25
40000/40000 - 13s - loss: 0.6782 - acc: 0.8624 - val_loss: 1.0224 - val_acc: 0.7660
Epoch 16/25
40000/40000 - 12s - loss: 0.6439 - acc: 0.8727 - val_loss: 1.0515 - val_acc: 0.7622
Epoch 17/25
40000/40000 - 12s - loss: 0.6134 - acc: 0.8838 - val_loss: 1.0397 - val_acc: 0.7640
Epoch 18/25
40000/40000 - 13s - loss: 0.5812 - acc: 0.8935 - val_loss: 1.0686 - val_acc: 0.7664
Epoch 19/25
40000/40000 - 13s - loss: 0.5579 - acc: 0.9033 - val_loss: 1.1098 - val_acc: 0.7611
Epoch 20/25
40000/40000 - 13s - loss: 0.5293 - acc: 0.9123 - val_loss: 1.1008 - val_acc: 0.7646
Epoch 21/25
40000/40000 - 12s - loss: 0.5046 - acc: 0.9221 - val_loss: 1.0843 - val_acc: 0.7717
Epoch 22/25
40000/40000 - 13s - loss: 0.4880 - acc: 0.9258 - val_loss: 1.1731 - val_acc: 0.7588
Epoch 23/25
40000/40000 - 12s - loss: 0.4678 - acc: 0.9329 - val_loss: 1.2208 - val_acc: 0.7593
Epoch 24/25
40000/40000 - 12s - loss: 0.4534 - acc: 0.9373 - val_loss: 1.2810 - val_acc: 0.7424
Epoch 25/25
40000/40000 - 13s - loss: 0.4402 - acc: 0.9418 - val_loss: 1.1663 - val_acc: 0.7663

Now, we can see that the training and validation loss are closer to each other! We are overfitting a bit less; however, we can improve this model. Let's continue to try with different weight decays.

Weight decay of 0.00005

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape,
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.00005)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dense(num_classes))
```

```

model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_2_part2.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 1.6239 - acc: 0.4171 - val_loss: 1.3006 - val_acc: 0.5461

Epoch 2/25

40000/40000 - 13s - loss: 1.2038 - acc: 0.5789 - val_loss: 1.1822 - val_acc: 0.6013

Epoch 3/25

40000/40000 - 13s - loss: 0.9981 - acc: 0.6535 - val_loss: 0.9808 - val_acc: 0.6609

Epoch 4/25

40000/40000 - 13s - loss: 0.8509 - acc: 0.7084 - val_loss: 0.8734 - val_acc: 0.7075

Epoch 5/25

40000/40000 - 13s - loss: 0.7345 - acc: 0.7496 - val_loss: 0.9499 - val_acc: 0.6783

Epoch 6/25

40000/40000 - 13s - loss: 0.6365 - acc: 0.7861 - val_loss: 0.8492 - val_acc: 0.7271

Epoch 7/25

40000/40000 - 13s - loss: 0.5523 - acc: 0.8152 - val_loss: 0.7895 - val_acc: 0.7444

Epoch 8/25

40000/40000 - 12s - loss: 0.4699 - acc: 0.8444 - val_loss: 0.8096 - val_acc: 0.7415

Epoch 9/25

40000/40000 - 12s - loss: 0.4002 - acc: 0.8667 - val_loss: 0.8677 - val_acc: 0.7435

Epoch 10/25
40000/40000 - 13s - loss: 0.3332 - acc: 0.8898 - val_loss: 0.8790 - val_acc: 0.7421

Epoch 11/25
40000/40000 - 12s - loss: 0.2740 - acc: 0.9112 - val_loss: 0.9327 - val_acc: 0.7487

Epoch 12/25
40000/40000 - 13s - loss: 0.2242 - acc: 0.9281 - val_loss: 1.0310 - val_acc: 0.7477

Epoch 13/25
40000/40000 - 13s - loss: 0.1843 - acc: 0.9405 - val_loss: 1.1074 - val_acc: 0.7431

Epoch 14/25
40000/40000 - 13s - loss: 0.1541 - acc: 0.9504 - val_loss: 1.2450 - val_acc: 0.7426

Epoch 15/25
40000/40000 - 12s - loss: 0.1338 - acc: 0.9588 - val_loss: 1.3136 - val_acc: 0.7441

Epoch 16/25
40000/40000 - 12s - loss: 0.1156 - acc: 0.9651 - val_loss: 1.4019 - val_acc: 0.7397

Epoch 17/25
40000/40000 - 13s - loss: 0.1079 - acc: 0.9685 - val_loss: 1.4684 - val_acc: 0.7474

Epoch 18/25
40000/40000 - 13s - loss: 0.0973 - acc: 0.9714 - val_loss: 1.4947 - val_acc: 0.7443

Epoch 19/25
40000/40000 - 13s - loss: 0.0894 - acc: 0.9748 - val_loss: 1.6713 - val_acc: 0.7397

Epoch 20/25
40000/40000 - 13s - loss: 0.0859 - acc: 0.9751 - val_loss: 1.6257 - val_acc: 0.7533

Epoch 21/25
40000/40000 - 12s - loss: 0.0828 - acc: 0.9778 - val_loss: 1.6815 - val_acc: 0.7401

Epoch 22/25
40000/40000 - 13s - loss: 0.0773 - acc: 0.9786 - val_loss: 1.9042 - val_acc: 0.7441

Epoch 23/25
40000/40000 - 13s - loss: 0.0749 - acc: 0.9802 - val_loss: 1.9526 - val_acc: 0.7429

Epoch 24/25
40000/40000 - 13s - loss: 0.0721 - acc: 0.9812 - val_loss: 1.9188 - val_acc: 0.7399

Epoch 25/25
40000/40000 - 13s - loss: 0.0729 - acc: 0.9810 - val_loss: 1.8911 - val_acc: 0.7455

Weight decay of 0.05

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape,
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0.05)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
```



```

model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history3 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_2_part3.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 4.7571 - acc: 0.3735 - val_loss: 2.5774 - val_acc: 0.4506

Epoch 2/25

40000/40000 - 13s - loss: 2.2108 - acc: 0.4658 - val_loss: 1.9172 - val_acc: 0.5055

Epoch 3/25

40000/40000 - 13s - loss: 1.8150 - acc: 0.5124 - val_loss: 1.6739 - val_acc: 0.5514

Epoch 4/25

40000/40000 - 13s - loss: 1.6372 - acc: 0.5481 - val_loss: 1.6642 - val_acc: 0.5279

Epoch 5/25

40000/40000 - 13s - loss: 1.5275 - acc: 0.5727 - val_loss: 1.6020 - val_acc: 0.5425

Epoch 6/25

40000/40000 - 13s - loss: 1.4457 - acc: 0.5938 - val_loss: 1.4015 - val_acc: 0.6060

Epoch 7/25

40000/40000 - 13s - loss: 1.3846 - acc: 0.6139 - val_loss: 1.4191 - val_acc: 0.6046

Epoch 8/25

40000/40000 - 13s - loss: 1.3289 - acc: 0.6331 - val_loss: 1.3470 - val_acc: 0.6224

Epoch 9/25

40000/40000 - 13s - loss: 1.2845 - acc: 0.6452 - val_loss: 1.3618 - val_acc: 0.6209

Epoch 10/25

40000/40000 - 13s - loss: 1.2481 - acc: 0.6583 - val_loss: 1.2855 - val_acc: 0.6403

Epoch 11/25

40000/40000 - 13s - loss: 1.2156 - acc: 0.6690 - val_loss: 1.2581 - val_acc: 0.6544
Epoch 12/25
40000/40000 - 13s - loss: 1.1878 - acc: 0.6782 - val_loss: 1.1940 - val_acc: 0.6797
Epoch 13/25
40000/40000 - 13s - loss: 1.1588 - acc: 0.6870 - val_loss: 1.2840 - val_acc: 0.6388
Epoch 14/25
40000/40000 - 13s - loss: 1.1346 - acc: 0.6952 - val_loss: 1.1559 - val_acc: 0.6885
Epoch 15/25
40000/40000 - 13s - loss: 1.1138 - acc: 0.7029 - val_loss: 1.1870 - val_acc: 0.6811
Epoch 16/25
40000/40000 - 13s - loss: 1.0946 - acc: 0.7104 - val_loss: 1.1661 - val_acc: 0.6839
Epoch 17/25
40000/40000 - 13s - loss: 1.0748 - acc: 0.7166 - val_loss: 1.1481 - val_acc: 0.6918
Epoch 18/25
40000/40000 - 13s - loss: 1.0567 - acc: 0.7224 - val_loss: 1.1470 - val_acc: 0.6939
Epoch 19/25
40000/40000 - 13s - loss: 1.0417 - acc: 0.7302 - val_loss: 1.1716 - val_acc: 0.6861
Epoch 20/25
40000/40000 - 13s - loss: 1.0221 - acc: 0.7346 - val_loss: 1.1957 - val_acc: 0.6837
Epoch 21/25
40000/40000 - 13s - loss: 1.0045 - acc: 0.7424 - val_loss: 1.1217 - val_acc: 0.7051
Epoch 22/25
40000/40000 - 13s - loss: 0.9901 - acc: 0.7479 - val_loss: 1.1151 - val_acc: 0.7088
Epoch 23/25
40000/40000 - 13s - loss: 0.9783 - acc: 0.7533 - val_loss: 1.1005 - val_acc: 0.7129
Epoch 24/25
40000/40000 - 13s - loss: 0.9622 - acc: 0.7575 - val_loss: 1.1166 - val_acc: 0.7120
Epoch 25/25
40000/40000 - 13s - loss: 0.9474 - acc: 0.7643 - val_loss: 1.1319 - val_acc: 0.7060

Weight decay of 0

```

[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape,
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(0)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',

```

```

        optimizer=opt,
        metrics=['accuracy'])

#Train model
history4 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_2_part4.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 1.5688 - acc: 0.4353 - val_loss: 1.3902 - val_acc: 0.5000

Epoch 2/25

40000/40000 - 12s - loss: 1.1755 - acc: 0.5842 - val_loss: 1.1016 - val_acc: 0.6126

Epoch 3/25

40000/40000 - 12s - loss: 0.9641 - acc: 0.6640 - val_loss: 0.9293 - val_acc: 0.6826

Epoch 4/25

40000/40000 - 12s - loss: 0.8164 - acc: 0.7182 - val_loss: 0.8757 - val_acc: 0.7034

Epoch 5/25

40000/40000 - 12s - loss: 0.7024 - acc: 0.7575 - val_loss: 0.8251 - val_acc: 0.7245

Epoch 6/25

40000/40000 - 12s - loss: 0.6080 - acc: 0.7913 - val_loss: 0.7696 - val_acc: 0.7414

Epoch 7/25

40000/40000 - 12s - loss: 0.5219 - acc: 0.8204 - val_loss: 0.7703 - val_acc: 0.7437

Epoch 8/25

40000/40000 - 12s - loss: 0.4447 - acc: 0.8470 - val_loss: 0.8055 - val_acc: 0.7478

Epoch 9/25

40000/40000 - 12s - loss: 0.3731 - acc: 0.8709 - val_loss: 0.9227 - val_acc: 0.7297

Epoch 10/25

40000/40000 - 12s - loss: 0.3073 - acc: 0.8944 - val_loss: 0.8496 - val_acc: 0.7526

Epoch 11/25

40000/40000 - 12s - loss: 0.2524 - acc: 0.9119 - val_loss: 0.8732 - val_acc:

0.7534
Epoch 12/25
40000/40000 - 12s - loss: 0.2036 - acc: 0.9294 - val_loss: 0.9980 - val_acc: 0.7461
Epoch 13/25
40000/40000 - 12s - loss: 0.1658 - acc: 0.9430 - val_loss: 1.0733 - val_acc: 0.7465
Epoch 14/25
40000/40000 - 12s - loss: 0.1387 - acc: 0.9516 - val_loss: 1.0687 - val_acc: 0.7581
Epoch 15/25
40000/40000 - 12s - loss: 0.1155 - acc: 0.9596 - val_loss: 1.3045 - val_acc: 0.7529
Epoch 16/25
40000/40000 - 12s - loss: 0.1000 - acc: 0.9648 - val_loss: 1.2257 - val_acc: 0.7524
Epoch 17/25
40000/40000 - 12s - loss: 0.0868 - acc: 0.9699 - val_loss: 1.4208 - val_acc: 0.7408
Epoch 18/25
40000/40000 - 12s - loss: 0.0778 - acc: 0.9721 - val_loss: 1.4997 - val_acc: 0.7478
Epoch 19/25
40000/40000 - 12s - loss: 0.0729 - acc: 0.9749 - val_loss: 1.5484 - val_acc: 0.7434
Epoch 20/25
40000/40000 - 12s - loss: 0.0673 - acc: 0.9764 - val_loss: 1.6378 - val_acc: 0.7366
Epoch 21/25
40000/40000 - 12s - loss: 0.0632 - acc: 0.9787 - val_loss: 1.5760 - val_acc: 0.7602
Epoch 22/25
40000/40000 - 12s - loss: 0.0571 - acc: 0.9801 - val_loss: 1.8291 - val_acc: 0.7553
Epoch 23/25
40000/40000 - 12s - loss: 0.0610 - acc: 0.9792 - val_loss: 2.0710 - val_acc: 0.7351
Epoch 24/25
40000/40000 - 12s - loss: 0.0545 - acc: 0.9812 - val_loss: 1.9122 - val_acc: 0.7548
Epoch 25/25
40000/40000 - 13s - loss: 0.0533 - acc: 0.9824 - val_loss: 1.8432 - val_acc: 0.7561

1.5.2 Validation


```
[0]: print("Weight decay default 0.0005")
plot_graphs(history1)


print("Weight decay default 0.00005")
plot_graphs(history2)

print("Weight decay default 0.05")
plot_graphs(history3)


print("Weight decay default 0")
plot_graphs(history4)
```


Weight decay default 0.0005







Weight decay default 0.00005







Weight decay default 0.05



Weight decay default 0



As we can see by the graphs above, weight decay of 0.05 does well to reduce the difference between the training and validation model loss. However, it sacrifices some accuracy. The weight decay of 0.0005 has a relatively good accuracy but has a large gap between the training and validation model loss (overfitting). We shall proceed with the 0.05 weight decay or combine the 0.0005 weight decay with another regularization to further prevent overfitting.

1.6 Model #3 - Regularization with Dropout Layer

We will try adding a dropout layer of 0.25 for each step. We will also try a dropout of 0.5 as well a model with increasing levels of dropout.

1.6.1 Dropout 0.25

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
```

```

model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dropout(0.25))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_3_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 1.7901 - acc: 0.3407 - val_loss: 1.5220 - val_acc: 0.4631

Epoch 2/25

40000/40000 - 13s - loss: 1.4017 - acc: 0.4910 - val_loss: 1.1977 - val_acc: 0.5840

Epoch 3/25

40000/40000 - 13s - loss: 1.2017 - acc: 0.5708 - val_loss: 1.0828 - val_acc: 0.6257

Epoch 4/25

40000/40000 - 13s - loss: 1.0708 - acc: 0.6209 - val_loss: 0.9911 - val_acc: 0.6521

Epoch 5/25

40000/40000 - 13s - loss: 0.9753 - acc: 0.6580 - val_loss: 0.8944 - val_acc: 0.6891

Epoch 6/25

40000/40000 - 13s - loss: 0.8979 - acc: 0.6852 - val_loss: 0.8760 - val_acc: 0.7004

Epoch 7/25

40000/40000 - 13s - loss: 0.8461 - acc: 0.7045 - val_loss: 0.8249 - val_acc: 0.7214
Epoch 8/25
40000/40000 - 13s - loss: 0.8025 - acc: 0.7218 - val_loss: 0.7604 - val_acc: 0.7410
Epoch 9/25
40000/40000 - 13s - loss: 0.7602 - acc: 0.7364 - val_loss: 0.7194 - val_acc: 0.7532
Epoch 10/25
40000/40000 - 13s - loss: 0.7308 - acc: 0.7475 - val_loss: 0.7207 - val_acc: 0.7509
Epoch 11/25
40000/40000 - 13s - loss: 0.6998 - acc: 0.7584 - val_loss: 0.6691 - val_acc: 0.7730
Epoch 12/25
40000/40000 - 13s - loss: 0.6758 - acc: 0.7655 - val_loss: 0.6590 - val_acc: 0.7732
Epoch 13/25
40000/40000 - 13s - loss: 0.6483 - acc: 0.7740 - val_loss: 0.6252 - val_acc: 0.7854
Epoch 14/25
40000/40000 - 13s - loss: 0.6339 - acc: 0.7804 - val_loss: 0.6528 - val_acc: 0.7794
Epoch 15/25
40000/40000 - 13s - loss: 0.6114 - acc: 0.7871 - val_loss: 0.6334 - val_acc: 0.7873
Epoch 16/25
40000/40000 - 13s - loss: 0.5966 - acc: 0.7945 - val_loss: 0.6137 - val_acc: 0.7920
Epoch 17/25
40000/40000 - 13s - loss: 0.5877 - acc: 0.7992 - val_loss: 0.6050 - val_acc: 0.7923
Epoch 18/25
40000/40000 - 13s - loss: 0.5687 - acc: 0.8051 - val_loss: 0.6249 - val_acc: 0.7930
Epoch 19/25
40000/40000 - 13s - loss: 0.5536 - acc: 0.8087 - val_loss: 0.6179 - val_acc: 0.7929
Epoch 20/25
40000/40000 - 13s - loss: 0.5498 - acc: 0.8114 - val_loss: 0.5928 - val_acc: 0.8014
Epoch 21/25
40000/40000 - 13s - loss: 0.5400 - acc: 0.8141 - val_loss: 0.5971 - val_acc: 0.8003
Epoch 22/25
40000/40000 - 13s - loss: 0.5340 - acc: 0.8179 - val_loss: 0.5911 - val_acc: 0.7992
Epoch 23/25

40000/40000 - 13s - loss: 0.5221 - acc: 0.8217 - val_loss: 0.6040 - val_acc: 0.7976
Epoch 24/25
40000/40000 - 13s - loss: 0.5218 - acc: 0.8241 - val_loss: 0.5686 - val_acc: 0.8103
Epoch 25/25
40000/40000 - 13s - loss: 0.5162 - acc: 0.8235 - val_loss: 0.6255 - val_acc: 0.7918

1.6.2 Dropout 0.5

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(256))
```

```

model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_3_part2.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 2.0588 - acc: 0.2271 - val_loss: 1.7303 - val_acc: 0.3730

Epoch 2/25

40000/40000 - 13s - loss: 1.6801 - acc: 0.3812 - val_loss: 1.6171 - val_acc: 0.4129

Epoch 3/25

40000/40000 - 13s - loss: 1.5185 - acc: 0.4439 - val_loss: 1.4503 - val_acc: 0.4793

Epoch 4/25

40000/40000 - 13s - loss: 1.4118 - acc: 0.4872 - val_loss: 1.2691 - val_acc: 0.5562

Epoch 5/25

40000/40000 - 13s - loss: 1.3326 - acc: 0.5207 - val_loss: 1.2628 - val_acc: 0.5499

Epoch 6/25

40000/40000 - 13s - loss: 1.2555 - acc: 0.5507 - val_loss: 1.1990 - val_acc: 0.5720

Epoch 7/25

40000/40000 - 13s - loss: 1.1948 - acc: 0.5756 - val_loss: 1.1330 - val_acc: 0.5991

Epoch 8/25

40000/40000 - 13s - loss: 1.1514 - acc: 0.5938 - val_loss: 1.0354 - val_acc: 0.6338

Epoch 9/25
40000/40000 - 13s - loss: 1.1173 - acc: 0.6089 - val_loss: 1.0001 - val_acc: 0.6501

Epoch 10/25
40000/40000 - 13s - loss: 1.0780 - acc: 0.6209 - val_loss: 1.0033 - val_acc: 0.6409

Epoch 11/25
40000/40000 - 13s - loss: 1.0516 - acc: 0.6295 - val_loss: 0.9677 - val_acc: 0.6583

Epoch 12/25
40000/40000 - 13s - loss: 1.0153 - acc: 0.6444 - val_loss: 0.9746 - val_acc: 0.6578

Epoch 13/25
40000/40000 - 13s - loss: 0.9984 - acc: 0.6492 - val_loss: 0.8731 - val_acc: 0.6953

Epoch 14/25
40000/40000 - 13s - loss: 0.9771 - acc: 0.6582 - val_loss: 0.9261 - val_acc: 0.6743

Epoch 15/25
40000/40000 - 13s - loss: 0.9666 - acc: 0.6617 - val_loss: 0.8705 - val_acc: 0.6893

Epoch 16/25
40000/40000 - 13s - loss: 0.9472 - acc: 0.6705 - val_loss: 0.8185 - val_acc: 0.7146

Epoch 17/25
40000/40000 - 13s - loss: 0.9359 - acc: 0.6771 - val_loss: 0.8238 - val_acc: 0.7174

Epoch 18/25
40000/40000 - 13s - loss: 0.9267 - acc: 0.6781 - val_loss: 0.8047 - val_acc: 0.7174

Epoch 19/25
40000/40000 - 13s - loss: 0.9149 - acc: 0.6809 - val_loss: 0.8146 - val_acc: 0.7218

Epoch 20/25
40000/40000 - 13s - loss: 0.9016 - acc: 0.6903 - val_loss: 0.8249 - val_acc: 0.7252

Epoch 21/25
40000/40000 - 13s - loss: 0.8949 - acc: 0.6902 - val_loss: 0.8127 - val_acc: 0.7278

Epoch 22/25
40000/40000 - 13s - loss: 0.8855 - acc: 0.6947 - val_loss: 0.8227 - val_acc: 0.7152

Epoch 23/25
40000/40000 - 13s - loss: 0.8846 - acc: 0.6941 - val_loss: 0.7816 - val_acc: 0.7389

Epoch 24/25
40000/40000 - 13s - loss: 0.8712 - acc: 0.7028 - val_loss: 0.7798 - val_acc: 0.7307

Epoch 25/25

40000/40000 - 13s - loss: 0.8723 - acc: 0.6981 - val_loss: 0.8437 - val_acc:
0.7122

1.6.3 Dropout increasing from 0.15 to 0.45

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

```

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history3 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_3_part3.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 20s - loss: 1.8625 - acc: 0.3076 - val_loss: 1.5225 - val_acc: 0.4586

Epoch 2/25

40000/40000 - 13s - loss: 1.4921 - acc: 0.4582 - val_loss: 1.3287 - val_acc: 0.5293

Epoch 3/25

40000/40000 - 13s - loss: 1.3044 - acc: 0.5333 - val_loss: 1.1574 - val_acc: 0.5823

Epoch 4/25

40000/40000 - 13s - loss: 1.1641 - acc: 0.5888 - val_loss: 1.0121 - val_acc: 0.6407

Epoch 5/25

40000/40000 - 13s - loss: 1.0642 - acc: 0.6251 - val_loss: 0.9195 - val_acc: 0.6829

Epoch 6/25

40000/40000 - 13s - loss: 0.9907 - acc: 0.6522 - val_loss: 0.8662 - val_acc: 0.7002

Epoch 7/25

40000/40000 - 13s - loss: 0.9320 - acc: 0.6741 - val_loss: 0.7970 - val_acc: 0.7180

Epoch 8/25

40000/40000 - 13s - loss: 0.8830 - acc: 0.6923 - val_loss: 0.7806 - val_acc: 0.7291

Epoch 9/25

40000/40000 - 13s - loss: 0.8439 - acc: 0.7070 - val_loss: 0.7376 - val_acc: 0.7429

Epoch 10/25

40000/40000 - 13s - loss: 0.8051 - acc: 0.7204 - val_loss: 0.7162 - val_acc:

0.7527
Epoch 11/25
40000/40000 - 13s - loss: 0.7763 - acc: 0.7310 - val_loss: 0.7041 - val_acc:
0.7589
Epoch 12/25
40000/40000 - 13s - loss: 0.7472 - acc: 0.7415 - val_loss: 0.7282 - val_acc:
0.7493
Epoch 13/25
40000/40000 - 13s - loss: 0.7234 - acc: 0.7516 - val_loss: 0.6809 - val_acc:
0.7650
Epoch 14/25
40000/40000 - 13s - loss: 0.7038 - acc: 0.7578 - val_loss: 0.6750 - val_acc:
0.7704
Epoch 15/25
40000/40000 - 13s - loss: 0.6895 - acc: 0.7637 - val_loss: 0.7180 - val_acc:
0.7520
Epoch 16/25
40000/40000 - 13s - loss: 0.6686 - acc: 0.7713 - val_loss: 0.6482 - val_acc:
0.7779
Epoch 17/25
40000/40000 - 13s - loss: 0.6540 - acc: 0.7771 - val_loss: 0.6316 - val_acc:
0.7806
Epoch 18/25
40000/40000 - 13s - loss: 0.6378 - acc: 0.7799 - val_loss: 0.6279 - val_acc:
0.7898
Epoch 19/25
40000/40000 - 13s - loss: 0.6244 - acc: 0.7867 - val_loss: 0.6171 - val_acc:
0.7898
Epoch 20/25
40000/40000 - 13s - loss: 0.6169 - acc: 0.7904 - val_loss: 0.6156 - val_acc:
0.7912
Epoch 21/25
40000/40000 - 13s - loss: 0.6081 - acc: 0.7905 - val_loss: 0.6084 - val_acc:
0.7913
Epoch 22/25
40000/40000 - 13s - loss: 0.5988 - acc: 0.7970 - val_loss: 0.6255 - val_acc:
0.7883
Epoch 23/25
40000/40000 - 13s - loss: 0.5910 - acc: 0.7969 - val_loss: 0.5970 - val_acc:
0.8003
Epoch 24/25
40000/40000 - 13s - loss: 0.5802 - acc: 0.8039 - val_loss: 0.6012 - val_acc:
0.7959
Epoch 25/25
40000/40000 - 13s - loss: 0.5768 - acc: 0.8035 - val_loss: 0.6313 - val_acc:
0.7891

1.6.4 Validation

The dropout layer of 0.25 and increasing from 0.15 to 0.45 seem to be better at both increasing accuracy of the training and validation set as well as reducing the loss between the two. Potentially increasing the number of epochs can smooth the curves (we will try increasing the number of epochs later).


We will either use these dropout layers for regularization to prevent overfitting or we will combine the dropout layer and weight decay together as weight decay alone is much worse at preventing overfitting than the dropout layer.


```
[0]: print("Dropout 0.25")
      plot_graphs(history1)

      print("Dropout 0.5")
      plot_graphs(history2)


      print("Dropout increasing from 0.15 -> 0.45")
      plot_graphs(history3)
```


Dropout 0.25







Dropout 0.5





Dropout increasing from 0.15 -> 0.45



1.7 Model #4 - Regularization Dropout and Weight Decay

Now, we will try combining the dropout and weight decay regularizations together to see if they are better than our best dropout layers (0.25 and 0.15->0.45)

1.7.1 Dropout 0.25, weight decay 0.0005

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(),
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
```



```

model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dropout(0.25))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_4_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 21s - loss: 3.1346 - acc: 0.3250 - val_loss: 2.3277 - val_acc: 0.4496

Epoch 2/25

40000/40000 - 13s - loss: 2.1222 - acc: 0.4586 - val_loss: 1.8253 - val_acc: 0.5261

Epoch 3/25

40000/40000 - 14s - loss: 1.7554 - acc: 0.5297 - val_loss: 1.5207 - val_acc: 0.5998

Epoch 4/25

40000/40000 - 13s - loss: 1.5361 - acc: 0.5784 - val_loss: 1.4004 - val_acc: 0.6203

Epoch 5/25

40000/40000 - 13s - loss: 1.3944 - acc: 0.6138 - val_loss: 1.2562 - val_acc: 0.6639

Epoch 6/25

40000/40000 - 13s - loss: 1.2991 - acc: 0.6358 - val_loss: 1.2189 - val_acc: 0.6656

Epoch 7/25

40000/40000 - 13s - loss: 1.2137 - acc: 0.6616 - val_loss: 1.2066 - val_acc: 0.6554

Epoch 8/25

40000/40000 - 13s - loss: 1.1523 - acc: 0.6775 - val_loss: 1.0730 - val_acc:

0.7088

Epoch 9/25
40000/40000 - 13s - loss: 1.1008 - acc: 0.6944 - val_loss: 1.0228 - val_acc: 0.7273

Epoch 10/25
40000/40000 - 13s - loss: 1.0529 - acc: 0.7084 - val_loss: 0.9856 - val_acc: 0.7329

Epoch 11/25
40000/40000 - 13s - loss: 1.0141 - acc: 0.7206 - val_loss: 0.9902 - val_acc: 0.7358

Epoch 12/25
40000/40000 - 13s - loss: 0.9839 - acc: 0.7287 - val_loss: 0.9479 - val_acc: 0.7408

Epoch 13/25
40000/40000 - 13s - loss: 0.9580 - acc: 0.7395 - val_loss: 0.8976 - val_acc: 0.7612

Epoch 14/25
40000/40000 - 14s - loss: 0.9316 - acc: 0.7456 - val_loss: 0.9037 - val_acc: 0.7542

Epoch 15/25
40000/40000 - 14s - loss: 0.9139 - acc: 0.7513 - val_loss: 0.9007 - val_acc: 0.7587

Epoch 16/25
40000/40000 - 13s - loss: 0.8906 - acc: 0.7563 - val_loss: 0.8689 - val_acc: 0.7693

Epoch 17/25
40000/40000 - 13s - loss: 0.8709 - acc: 0.7672 - val_loss: 0.8576 - val_acc: 0.7728

Epoch 18/25
40000/40000 - 13s - loss: 0.8589 - acc: 0.7696 - val_loss: 0.9687 - val_acc: 0.7350

Epoch 19/25
40000/40000 - 13s - loss: 0.8386 - acc: 0.7785 - val_loss: 0.8098 - val_acc: 0.7900

Epoch 20/25
40000/40000 - 14s - loss: 0.8283 - acc: 0.7777 - val_loss: 0.8056 - val_acc: 0.7880

Epoch 21/25
40000/40000 - 13s - loss: 0.8172 - acc: 0.7827 - val_loss: 0.8689 - val_acc: 0.7730

Epoch 22/25
40000/40000 - 13s - loss: 0.8071 - acc: 0.7879 - val_loss: 0.8140 - val_acc: 0.7878

Epoch 23/25
40000/40000 - 13s - loss: 0.7980 - acc: 0.7911 - val_loss: 0.7990 - val_acc: 0.7913

Epoch 24/25
40000/40000 - 13s - loss: 0.7880 - acc: 0.7925 - val_loss: 0.8101 - val_acc:

0.7913
Epoch 25/25
40000/40000 - 13s - loss: 0.7763 - acc: 0.7996 - val_loss: 0.7612 - val_acc:
0.8045

1.7.2 Dropout increasing, weight decay 0.0005

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(),
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))
```

```

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_4_part2.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 21s - loss: 3.1879 - acc: 0.3022 - val_loss: 2.6121 - val_acc: 0.3612

Epoch 2/25

40000/40000 - 13s - loss: 2.1792 - acc: 0.4294 - val_loss: 1.8680 - val_acc: 0.5020

Epoch 3/25

40000/40000 - 13s - loss: 1.8000 - acc: 0.5050 - val_loss: 1.6364 - val_acc: 0.5520

Epoch 4/25

40000/40000 - 13s - loss: 1.5745 - acc: 0.5561 - val_loss: 1.4080 - val_acc: 0.6173

Epoch 5/25

40000/40000 - 13s - loss: 1.4192 - acc: 0.6004 - val_loss: 1.2522 - val_acc: 0.6520

Epoch 6/25

40000/40000 - 13s - loss: 1.3133 - acc: 0.6307 - val_loss: 1.1678 - val_acc: 0.6755

Epoch 7/25

40000/40000 - 13s - loss: 1.2342 - acc: 0.6541 - val_loss: 1.1183 - val_acc: 0.6881

Epoch 8/25

40000/40000 - 13s - loss: 1.1647 - acc: 0.6741 - val_loss: 1.0753 - val_acc: 0.6958
Epoch 9/25
40000/40000 - 13s - loss: 1.1156 - acc: 0.6873 - val_loss: 1.0295 - val_acc: 0.7182
Epoch 10/25
40000/40000 - 13s - loss: 1.0747 - acc: 0.6982 - val_loss: 1.0007 - val_acc: 0.7243
Epoch 11/25
40000/40000 - 13s - loss: 1.0451 - acc: 0.7082 - val_loss: 0.9828 - val_acc: 0.7330
Epoch 12/25
40000/40000 - 13s - loss: 1.0151 - acc: 0.7196 - val_loss: 0.9283 - val_acc: 0.7488
Epoch 13/25
40000/40000 - 13s - loss: 0.9891 - acc: 0.7254 - val_loss: 0.9379 - val_acc: 0.7489
Epoch 14/25
40000/40000 - 13s - loss: 0.9642 - acc: 0.7332 - val_loss: 0.9045 - val_acc: 0.7576
Epoch 15/25
40000/40000 - 13s - loss: 0.9430 - acc: 0.7432 - val_loss: 0.9185 - val_acc: 0.7523
Epoch 16/25
40000/40000 - 13s - loss: 0.9294 - acc: 0.7466 - val_loss: 0.8538 - val_acc: 0.7718
Epoch 17/25
40000/40000 - 13s - loss: 0.9123 - acc: 0.7518 - val_loss: 0.8556 - val_acc: 0.7765
Epoch 18/25
40000/40000 - 13s - loss: 0.8966 - acc: 0.7582 - val_loss: 0.8632 - val_acc: 0.7700
Epoch 19/25
40000/40000 - 13s - loss: 0.8817 - acc: 0.7622 - val_loss: 0.8311 - val_acc: 0.7802
Epoch 20/25
40000/40000 - 13s - loss: 0.8690 - acc: 0.7673 - val_loss: 0.8142 - val_acc: 0.7868
Epoch 21/25
40000/40000 - 13s - loss: 0.8574 - acc: 0.7715 - val_loss: 0.8266 - val_acc: 0.7846
Epoch 22/25
40000/40000 - 13s - loss: 0.8488 - acc: 0.7732 - val_loss: 0.8306 - val_acc: 0.7810
Epoch 23/25
40000/40000 - 13s - loss: 0.8364 - acc: 0.7796 - val_loss: 0.8020 - val_acc: 0.7907
Epoch 24/25

40000/40000 - 13s - loss: 0.8237 - acc: 0.7841 - val_loss: 0.8122 - val_acc: 0.7862

Epoch 25/25

40000/40000 - 13s - loss: 0.8182 - acc: 0.7853 - val_loss: 0.7901 - val_acc: 0.7937

1.7.3 Dropout increasing, weight decay 0.0005, both alternating

```
[0]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(),
                input_shape=shape))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
```

```

model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history3 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_4_part3.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 21s - loss: 2.8603 - acc: 0.3458 - val_loss: 2.1650 - val_acc: 0.4605

Epoch 2/25

40000/40000 - 13s - loss: 1.9755 - acc: 0.4733 - val_loss: 1.7810 - val_acc: 0.5016

Epoch 3/25

40000/40000 - 13s - loss: 1.6072 - acc: 0.5516 - val_loss: 1.4278 - val_acc: 0.6018

Epoch 4/25

40000/40000 - 13s - loss: 1.3911 - acc: 0.6022 - val_loss: 1.2760 - val_acc: 0.6311

Epoch 5/25

40000/40000 - 13s - loss: 1.2491 - acc: 0.6403 - val_loss: 1.1161 - val_acc: 0.6747

Epoch 6/25

40000/40000 - 13s - loss: 1.1450 - acc: 0.6676 - val_loss: 1.0785 - val_acc: 0.6857

Epoch 7/25

40000/40000 - 13s - loss: 1.0644 - acc: 0.6915 - val_loss: 0.9764 - val_acc: 0.7166

Epoch 8/25

40000/40000 - 13s - loss: 0.9969 - acc: 0.7087 - val_loss: 1.0232 - val_acc: 0.7016

Epoch 9/25

40000/40000 - 13s - loss: 0.9486 - acc: 0.7249 - val_loss: 0.9733 - val_acc: 0.7167
Epoch 10/25
40000/40000 - 13s - loss: 0.9062 - acc: 0.7387 - val_loss: 0.9004 - val_acc: 0.7472
Epoch 11/25
40000/40000 - 13s - loss: 0.8656 - acc: 0.7526 - val_loss: 0.8845 - val_acc: 0.7377
Epoch 12/25
40000/40000 - 13s - loss: 0.8399 - acc: 0.7622 - val_loss: 0.8427 - val_acc: 0.7606
Epoch 13/25
40000/40000 - 13s - loss: 0.8040 - acc: 0.7715 - val_loss: 0.7979 - val_acc: 0.7722
Epoch 14/25
40000/40000 - 13s - loss: 0.7818 - acc: 0.7802 - val_loss: 0.8128 - val_acc: 0.7686
Epoch 15/25
40000/40000 - 13s - loss: 0.7593 - acc: 0.7896 - val_loss: 0.9388 - val_acc: 0.7376
Epoch 16/25
40000/40000 - 13s - loss: 0.7350 - acc: 0.7959 - val_loss: 0.7680 - val_acc: 0.7836
Epoch 17/25
40000/40000 - 13s - loss: 0.7188 - acc: 0.8031 - val_loss: 0.8194 - val_acc: 0.7727
Epoch 18/25
40000/40000 - 13s - loss: 0.7043 - acc: 0.8053 - val_loss: 0.7605 - val_acc: 0.7919
Epoch 19/25
40000/40000 - 13s - loss: 0.6919 - acc: 0.8120 - val_loss: 0.7407 - val_acc: 0.7986
Epoch 20/25
40000/40000 - 13s - loss: 0.6725 - acc: 0.8203 - val_loss: 0.7495 - val_acc: 0.7943
Epoch 21/25
40000/40000 - 13s - loss: 0.6624 - acc: 0.8230 - val_loss: 0.7800 - val_acc: 0.7881
Epoch 22/25
40000/40000 - 13s - loss: 0.6572 - acc: 0.8255 - val_loss: 0.8197 - val_acc: 0.7781
Epoch 23/25
40000/40000 - 13s - loss: 0.6480 - acc: 0.8292 - val_loss: 0.7262 - val_acc: 0.8064
Epoch 24/25
40000/40000 - 13s - loss: 0.6385 - acc: 0.8342 - val_loss: 0.7461 - val_acc: 0.8023
Epoch 25/25

40000/40000 - 13s - loss: 0.6316 - acc: 0.8367 - val_loss: 0.7916 - val_acc: 0.8032

1.7.4 Validation

From the graphs below, it seems like alternating increasing dropout from 0.15 to 0.45 and and weight decay of 0.0005 increases accuracy and minimizes the training and validation loss the most. However, just increasing the dropout 0.15 to 0.45 produced similar accuracies but minimized the training and validation loss moreso than the combination of dropout and weight decay.

For the section below, we will proceed by using seeing how normalization impacts our best models so far.


1. increasing dropout from 0.15 to 0.45
2. increasing dropout from 0.15 to 0.45 and and weight decay of 0.0005)


```
[0]: print("Dropout 0.25, weight decay 0.0005")
      plot_graphs(history1)

      print("Dropout increasing from 0.15 -> 0.45, weight decay 0.0005")
      plot_graphs(history2)



      print("Alternating dropout (increasing) and weight decay 0.0005")
      plot_graphs(history3)
```

Dropout 0.25, weight decay 0.0005







Dropout increasing from 0.15 -> 0.45, weight decay 0.0005



Alternating dropout (increasing) and weight decay 0.0005



1.8 Model #5 - Normalization

Now that we added regularization to the model, let's also add normalization. We will add Batch Normalization, which will normalize the activations of the previous hidden layer at each batch.

Let's see how regularizing and normalizing the layers compares to just regularizing.

1.8.1 Dropout increasing with normalization

```
[15]: from tensorflow.keras.layers import BatchNormalization
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
```

```

        kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_5_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 12s - loss: 2.2320 - acc: 0.3167 - val_loss: 1.5173 - val_acc: 0.4633

Epoch 2/25

40000/40000 - 11s - loss: 1.6244 - acc: 0.4555 - val_loss: 1.2026 - val_acc: 0.5707

Epoch 3/25

40000/40000 - 11s - loss: 1.3475 - acc: 0.5423 - val_loss: 1.0285 - val_acc: 0.6400

Epoch 4/25

40000/40000 - 11s - loss: 1.1582 - acc: 0.5988 - val_loss: 1.0622 - val_acc: 0.6421

Epoch 5/25


40000/40000 - 11s - loss: 1.0363 - acc: 0.6397 - val_loss: 0.9257 - val_acc: 0.6761


Epoch 6/25
40000/40000 - 11s - loss: 0.9441 - acc: 0.6712 - val_loss: 0.7996 - val_acc: 0.7215
Epoch 7/25
40000/40000 - 11s - loss: 0.8747 - acc: 0.6949 - val_loss: 0.7942 - val_acc: 0.7193
Epoch 8/25
40000/40000 - 11s - loss: 0.8148 - acc: 0.7177 - val_loss: 0.7697 - val_acc: 0.7230
Epoch 9/25
40000/40000 - 11s - loss: 0.7700 - acc: 0.7346 - val_loss: 0.7907 - val_acc: 0.7294
Epoch 10/25
40000/40000 - 11s - loss: 0.7267 - acc: 0.7487 - val_loss: 0.7276 - val_acc: 0.7484
Epoch 11/25
40000/40000 - 11s - loss: 0.6934 - acc: 0.7591 - val_loss: 0.6376 - val_acc: 0.7796
Epoch 12/25
40000/40000 - 11s - loss: 0.6618 - acc: 0.7711 - val_loss: 0.6953 - val_acc: 0.7604
Epoch 13/25
40000/40000 - 11s - loss: 0.6405 - acc: 0.7786 - val_loss: 0.6269 - val_acc: 0.7778
Epoch 14/25
40000/40000 - 11s - loss: 0.6096 - acc: 0.7897 - val_loss: 0.6706 - val_acc: 0.7675
Epoch 15/25
40000/40000 - 11s - loss: 0.5910 - acc: 0.7960 - val_loss: 0.5888 - val_acc: 0.7935
Epoch 16/25
40000/40000 - 11s - loss: 0.5651 - acc: 0.8027 - val_loss: 0.5636 - val_acc: 0.8073
Epoch 17/25
40000/40000 - 11s - loss: 0.5439 - acc: 0.8132 - val_loss: 0.5399 - val_acc: 0.8116
Epoch 18/25
40000/40000 - 11s - loss: 0.5281 - acc: 0.8179 - val_loss: 0.5330 - val_acc: 0.8156
Epoch 19/25
40000/40000 - 11s - loss: 0.5100 - acc: 0.8241 - val_loss: 0.5447 - val_acc: 0.8100
Epoch 20/25
40000/40000 - 11s - loss: 0.4936 - acc: 0.8315 - val_loss: 0.5456 - val_acc: 0.8158
Epoch 21/25
40000/40000 - 11s - loss: 0.4713 - acc: 0.8369 - val_loss: 0.4975 - val_acc: 0.8298

```
Epoch 22/25  
40000/40000 - 11s - loss: 0.4656 - acc: 0.8394 - val_loss: 0.5683 - val_acc:  
0.8054  
Epoch 23/25  
40000/40000 - 11s - loss: 0.4404 - acc: 0.8496 - val_loss: 0.5148 - val_acc:  
0.8253  
Epoch 24/25  
40000/40000 - 11s - loss: 0.4306 - acc: 0.8505 - val_loss: 0.5101 - val_acc:  
0.8321  
Epoch 25/25  
40000/40000 - 11s - loss: 0.4206 - acc: 0.8530 - val_loss: 0.5218 - val_acc:  
0.8290
```

```
[59]: print("Dropout increasing with normalization")  
      plot_graphs(history1)
```

Dropout increasing with normalization





We see that the model above is overfitting the training data. We will add more regularization (both dropout and weight decay).

Dropout increasing 0.15 -> 0.45 and weight decay = 0.005 with normalization

```
[22]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2(),
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
```

```

        kernel_initializer='random_uniform',
        kernel_regularizer=regularizers.l2())
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform',
                kernel_regularizer=regularizers.l2()))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,

```

```
validation_split=0.2,  
shuffle=True)
```

```
#Save model
```

```
model.save('model_5_part2.h5', overwrite=True)
```

Train on 40000 samples, validate on 10000 samples

Epoch 1/25

40000/40000 - 13s - loss: 4.5092 - acc: 0.3197 - val_loss: 3.5261 - val_acc: 0.5139

Epoch 2/25

40000/40000 - 12s - loss: 3.5274 - acc: 0.4775 - val_loss: 3.0077 - val_acc: 0.5600

Epoch 3/25

40000/40000 - 12s - loss: 2.8864 - acc: 0.5605 - val_loss: 2.4447 - val_acc: 0.6447

Epoch 4/25

40000/40000 - 12s - loss: 2.4109 - acc: 0.6084 - val_loss: 2.0729 - val_acc: 0.6781

Epoch 5/25

40000/40000 - 12s - loss: 2.0308 - acc: 0.6532 - val_loss: 1.7648 - val_acc: 0.7029

Epoch 6/25

40000/40000 - 12s - loss: 1.7419 - acc: 0.6842 - val_loss: 1.5506 - val_acc: 0.7170

Epoch 7/25

40000/40000 - 12s - loss: 1.5166 - acc: 0.7105 - val_loss: 1.3802 - val_acc: 0.7383

Epoch 8/25

40000/40000 - 12s - loss: 1.3595 - acc: 0.7301 - val_loss: 1.2005 - val_acc: 0.7689

Epoch 9/25

40000/40000 - 12s - loss: 1.2229 - acc: 0.7485 - val_loss: 1.1307 - val_acc: 0.7699

Epoch 10/25

40000/40000 - 12s - loss: 1.1216 - acc: 0.7602 - val_loss: 1.0203 - val_acc: 0.7881

Epoch 11/25

40000/40000 - 12s - loss: 1.0398 - acc: 0.7743 - val_loss: 0.9515 - val_acc: 0.7976

Epoch 12/25

40000/40000 - 12s - loss: 0.9682 - acc: 0.7852 - val_loss: 0.9465 - val_acc: 0.7887

Epoch 13/25

40000/40000 - 12s - loss: 0.9151 - acc: 0.7957 - val_loss: 0.8735 - val_acc: 0.8033

Epoch 14/25

```

40000/40000 - 12s - loss: 0.8615 - acc: 0.8084 - val_loss: 0.8416 - val_acc:
0.8124
Epoch 15/25
40000/40000 - 12s - loss: 0.8267 - acc: 0.8134 - val_loss: 0.8219 - val_acc:
0.8157
Epoch 16/25
40000/40000 - 12s - loss: 0.7922 - acc: 0.8224 - val_loss: 0.8233 - val_acc:
0.8110
Epoch 17/25
40000/40000 - 12s - loss: 0.7622 - acc: 0.8281 - val_loss: 0.7941 - val_acc:
0.8175
Epoch 18/25
40000/40000 - 12s - loss: 0.7429 - acc: 0.8324 - val_loss: 0.7884 - val_acc:
0.8180
Epoch 19/25
40000/40000 - 12s - loss: 0.7221 - acc: 0.8374 - val_loss: 0.7691 - val_acc:
0.8247
Epoch 20/25
40000/40000 - 12s - loss: 0.7021 - acc: 0.8443 - val_loss: 0.7579 - val_acc:
0.8250
Epoch 21/25
40000/40000 - 12s - loss: 0.6784 - acc: 0.8485 - val_loss: 0.7434 - val_acc:
0.8312
Epoch 22/25
40000/40000 - 12s - loss: 0.6614 - acc: 0.8535 - val_loss: 0.8350 - val_acc:
0.8030
Epoch 23/25
40000/40000 - 12s - loss: 0.6556 - acc: 0.8569 - val_loss: 0.7741 - val_acc:
0.8262
Epoch 24/25
40000/40000 - 12s - loss: 0.6423 - acc: 0.8606 - val_loss: 0.7813 - val_acc:
0.8192
Epoch 25/25
40000/40000 - 12s - loss: 0.6329 - acc: 0.8626 - val_loss: 0.7140 - val_acc:
0.8414



```

```

[23]: print("Dropout increasing from 0.15 -> 0.45, weight decay 0.0005, normalization_
      ↪after each layer")
      plot_graphs(history2)

```

Dropout increasing from 0.15 -> 0.45, weight decay 0.0005, normalization after each layer



1.8.2 Validation

Looking at the graphs above as well as the graphs in our model #3 and model #4, normalization has improved both our accuracy and has reduced the validation/training loss. We will keep this in our model. We will also use the dropout increasing for the regularization because it reduces overfitting more than if we add weight decay.

1.9 Model #6 - Optimizers

In the first model, we arbitrarily chose the “RMSprop” optimizer. Now we will try several other optimizers to see which one produces the best results. However, we will also increase the number of epochs.

1.9.1 RMSprop

```
[0]: epochs = 25

[27]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
```

```

model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_6_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/50

40000/40000 - 13s - loss: 2.2848 - acc: 0.3016 - val_loss: 1.5005 - val_acc: 0.4640

Epoch 2/50

40000/40000 - 12s - loss: 1.6665 - acc: 0.4458 - val_loss: 1.3384 - val_acc: 0.5169

Epoch 3/50

40000/40000 - 12s - loss: 1.4028 - acc: 0.5221 - val_loss: 1.0763 - val_acc: 0.6161

Epoch 4/50

40000/40000 - 12s - loss: 1.2039 - acc: 0.5833 - val_loss: 1.0893 - val_acc: 0.6208

Epoch 5/50
40000/40000 - 12s - loss: 1.0739 - acc: 0.6280 - val_loss: 0.8896 - val_acc: 0.6819
Epoch 6/50
40000/40000 - 12s - loss: 0.9801 - acc: 0.6603 - val_loss: 0.8239 - val_acc: 0.7085
Epoch 7/50
40000/40000 - 12s - loss: 0.8989 - acc: 0.6848 - val_loss: 0.7973 - val_acc: 0.7216
Epoch 8/50
40000/40000 - 12s - loss: 0.8389 - acc: 0.7101 - val_loss: 0.7316 - val_acc: 0.7485
Epoch 9/50
40000/40000 - 12s - loss: 0.7892 - acc: 0.7254 - val_loss: 0.6843 - val_acc: 0.7588
Epoch 10/50
40000/40000 - 12s - loss: 0.7480 - acc: 0.7391 - val_loss: 0.7494 - val_acc: 0.7380
Epoch 11/50
40000/40000 - 12s - loss: 0.7120 - acc: 0.7510 - val_loss: 0.7443 - val_acc: 0.7381
Epoch 12/50
40000/40000 - 12s - loss: 0.6822 - acc: 0.7612 - val_loss: 0.6649 - val_acc: 0.7671
Epoch 13/50
40000/40000 - 11s - loss: 0.6493 - acc: 0.7737 - val_loss: 0.6046 - val_acc: 0.7851
Epoch 14/50
40000/40000 - 12s - loss: 0.6259 - acc: 0.7837 - val_loss: 0.5968 - val_acc: 0.7922
Epoch 15/50
40000/40000 - 11s - loss: 0.5990 - acc: 0.7928 - val_loss: 0.5811 - val_acc: 0.7962
Epoch 16/50
40000/40000 - 11s - loss: 0.5817 - acc: 0.7970 - val_loss: 0.5651 - val_acc: 0.8062
Epoch 17/50
40000/40000 - 12s - loss: 0.5573 - acc: 0.8062 - val_loss: 0.5762 - val_acc: 0.8045
Epoch 18/50
40000/40000 - 12s - loss: 0.5341 - acc: 0.8132 - val_loss: 0.5574 - val_acc: 0.8094
Epoch 19/50
40000/40000 - 12s - loss: 0.5200 - acc: 0.8192 - val_loss: 0.5734 - val_acc: 0.8012
Epoch 20/50
40000/40000 - 12s - loss: 0.5065 - acc: 0.8217 - val_loss: 0.5341 - val_acc: 0.8174

Epoch 21/50
40000/40000 - 12s - loss: 0.4886 - acc: 0.8320 - val_loss: 0.5383 - val_acc: 0.8165
Epoch 22/50
40000/40000 - 12s - loss: 0.4693 - acc: 0.8363 - val_loss: 0.5446 - val_acc: 0.8134
Epoch 23/50
40000/40000 - 12s - loss: 0.4551 - acc: 0.8422 - val_loss: 0.5428 - val_acc: 0.8164
Epoch 24/50
40000/40000 - 11s - loss: 0.4403 - acc: 0.8488 - val_loss: 0.5083 - val_acc: 0.8256
Epoch 25/50
40000/40000 - 12s - loss: 0.4250 - acc: 0.8525 - val_loss: 0.5178 - val_acc: 0.8239
Epoch 26/50
40000/40000 - 11s - loss: 0.4153 - acc: 0.8560 - val_loss: 0.5178 - val_acc: 0.8239
Epoch 27/50
40000/40000 - 12s - loss: 0.4023 - acc: 0.8578 - val_loss: 0.4998 - val_acc: 0.8312
Epoch 28/50
40000/40000 - 12s - loss: 0.3926 - acc: 0.8638 - val_loss: 0.5105 - val_acc: 0.8258
Epoch 29/50
40000/40000 - 11s - loss: 0.3823 - acc: 0.8662 - val_loss: 0.5172 - val_acc: 0.8263
Epoch 30/50
40000/40000 - 11s - loss: 0.3721 - acc: 0.8708 - val_loss: 0.5156 - val_acc: 0.8324
Epoch 31/50
40000/40000 - 12s - loss: 0.3631 - acc: 0.8730 - val_loss: 0.5158 - val_acc: 0.8291
Epoch 32/50
40000/40000 - 12s - loss: 0.3537 - acc: 0.8744 - val_loss: 0.5063 - val_acc: 0.8332
Epoch 33/50
40000/40000 - 12s - loss: 0.3429 - acc: 0.8781 - val_loss: 0.5079 - val_acc: 0.8354
Epoch 34/50
40000/40000 - 11s - loss: 0.3362 - acc: 0.8820 - val_loss: 0.5372 - val_acc: 0.8295
Epoch 35/50
40000/40000 - 11s - loss: 0.3273 - acc: 0.8868 - val_loss: 0.5313 - val_acc: 0.8329
Epoch 36/50
40000/40000 - 11s - loss: 0.3194 - acc: 0.8887 - val_loss: 0.5354 - val_acc: 0.8301

Epoch 37/50
40000/40000 - 11s - loss: 0.3119 - acc: 0.8922 - val_loss: 0.5313 - val_acc: 0.8284
Epoch 38/50
40000/40000 - 11s - loss: 0.3087 - acc: 0.8910 - val_loss: 0.5158 - val_acc: 0.8371
Epoch 39/50
40000/40000 - 11s - loss: 0.2958 - acc: 0.8959 - val_loss: 0.4947 - val_acc: 0.8403
Epoch 40/50
40000/40000 - 11s - loss: 0.2930 - acc: 0.8963 - val_loss: 0.5207 - val_acc: 0.8371
Epoch 41/50
40000/40000 - 11s - loss: 0.2902 - acc: 0.8985 - val_loss: 0.5481 - val_acc: 0.8279
Epoch 42/50
40000/40000 - 11s - loss: 0.2842 - acc: 0.9000 - val_loss: 0.5348 - val_acc: 0.8306
Epoch 43/50
40000/40000 - 11s - loss: 0.2786 - acc: 0.9022 - val_loss: 0.5022 - val_acc: 0.8409
Epoch 44/50
40000/40000 - 11s - loss: 0.2714 - acc: 0.9044 - val_loss: 0.5028 - val_acc: 0.8440
Epoch 45/50
40000/40000 - 11s - loss: 0.2627 - acc: 0.9070 - val_loss: 0.5343 - val_acc: 0.8343
Epoch 46/50
40000/40000 - 11s - loss: 0.2592 - acc: 0.9076 - val_loss: 0.5062 - val_acc: 0.8438
Epoch 47/50
40000/40000 - 11s - loss: 0.2552 - acc: 0.9105 - val_loss: 0.5159 - val_acc: 0.8442
Epoch 48/50
40000/40000 - 11s - loss: 0.2467 - acc: 0.9124 - val_loss: 0.5299 - val_acc: 0.8397
Epoch 49/50
40000/40000 - 11s - loss: 0.2438 - acc: 0.9139 - val_loss: 0.5036 - val_acc: 0.8468
Epoch 50/50
40000/40000 - 11s - loss: 0.2410 - acc: 0.9136 - val_loss: 0.5403 - val_acc: 0.8386

1.9.2 SGD - Stochastic gradient descent optimizer

```
[32]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
```

```

model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.SGD(learning_rate=0.01, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_6_part2.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/50

40000/40000 - 14s - loss: 1.8051 - acc: 0.3739 - val_loss: 1.4150 - val_acc: 0.5021

Epoch 2/50

40000/40000 - 11s - loss: 1.3103 - acc: 0.5346 - val_loss: 1.0651 - val_acc: 0.6146

Epoch 3/50

40000/40000 - 11s - loss: 1.0742 - acc: 0.6229 - val_loss: 0.9073 - val_acc: 0.6793

Epoch 4/50

40000/40000 - 11s - loss: 0.9253 - acc: 0.6745 - val_loss: 0.7608 - val_acc: 0.7305

Epoch 5/50

40000/40000 - 11s - loss: 0.8235 - acc: 0.7096 - val_loss: 0.6986 - val_acc: 0.7562

Epoch 6/50

40000/40000 - 11s - loss: 0.7481 - acc: 0.7397 - val_loss: 0.6321 - val_acc: 0.7782

Epoch 7/50

40000/40000 - 11s - loss: 0.6873 - acc: 0.7604 - val_loss: 0.6395 - val_acc: 0.7756

Epoch 8/50

40000/40000 - 11s - loss: 0.6429 - acc: 0.7775 - val_loss: 0.5805 - val_acc: 0.8006

Epoch 9/50

40000/40000 - 11s - loss: 0.6040 - acc: 0.7907 - val_loss: 0.5761 - val_acc:

0.8011
Epoch 10/50
40000/40000 - 11s - loss: 0.5760 - acc: 0.8007 - val_loss: 0.5828 - val_acc: 0.7995
Epoch 11/50
40000/40000 - 11s - loss: 0.5420 - acc: 0.8122 - val_loss: 0.5641 - val_acc: 0.8053
Epoch 12/50
40000/40000 - 11s - loss: 0.5152 - acc: 0.8198 - val_loss: 0.5771 - val_acc: 0.8036
Epoch 13/50
40000/40000 - 11s - loss: 0.4873 - acc: 0.8316 - val_loss: 0.5179 - val_acc: 0.8218
Epoch 14/50
40000/40000 - 11s - loss: 0.4650 - acc: 0.8379 - val_loss: 0.5403 - val_acc: 0.8156
Epoch 15/50
40000/40000 - 11s - loss: 0.4489 - acc: 0.8433 - val_loss: 0.5159 - val_acc: 0.8278
Epoch 16/50
40000/40000 - 11s - loss: 0.4346 - acc: 0.8487 - val_loss: 0.5162 - val_acc: 0.8297
Epoch 17/50
40000/40000 - 11s - loss: 0.4115 - acc: 0.8559 - val_loss: 0.4904 - val_acc: 0.8379
Epoch 18/50
40000/40000 - 11s - loss: 0.4002 - acc: 0.8603 - val_loss: 0.5236 - val_acc: 0.8275
Epoch 19/50
40000/40000 - 11s - loss: 0.3815 - acc: 0.8654 - val_loss: 0.4962 - val_acc: 0.8407
Epoch 20/50
40000/40000 - 11s - loss: 0.3722 - acc: 0.8704 - val_loss: 0.4971 - val_acc: 0.8368
Epoch 21/50
40000/40000 - 11s - loss: 0.3562 - acc: 0.8753 - val_loss: 0.5031 - val_acc: 0.8358
Epoch 22/50
40000/40000 - 11s - loss: 0.3445 - acc: 0.8787 - val_loss: 0.5070 - val_acc: 0.8395
Epoch 23/50
40000/40000 - 11s - loss: 0.3322 - acc: 0.8841 - val_loss: 0.4973 - val_acc: 0.8416
Epoch 24/50
40000/40000 - 11s - loss: 0.3246 - acc: 0.8857 - val_loss: 0.5115 - val_acc: 0.8385
Epoch 25/50
40000/40000 - 11s - loss: 0.3080 - acc: 0.8924 - val_loss: 0.4983 - val_acc:

0.8450
Epoch 26/50
40000/40000 - 11s - loss: 0.3045 - acc: 0.8927 - val_loss: 0.4913 - val_acc: 0.8465
Epoch 27/50
40000/40000 - 11s - loss: 0.2944 - acc: 0.8961 - val_loss: 0.5057 - val_acc: 0.8419
Epoch 28/50
40000/40000 - 11s - loss: 0.2881 - acc: 0.8982 - val_loss: 0.5035 - val_acc: 0.8393
Epoch 29/50
40000/40000 - 11s - loss: 0.2751 - acc: 0.9019 - val_loss: 0.5169 - val_acc: 0.8406
Epoch 30/50
40000/40000 - 11s - loss: 0.2797 - acc: 0.9007 - val_loss: 0.5031 - val_acc: 0.8456
Epoch 31/50
40000/40000 - 11s - loss: 0.2511 - acc: 0.9115 - val_loss: 0.5202 - val_acc: 0.8424
Epoch 32/50
40000/40000 - 11s - loss: 0.2589 - acc: 0.9068 - val_loss: 0.5153 - val_acc: 0.8435
Epoch 33/50
40000/40000 - 11s - loss: 0.2536 - acc: 0.9103 - val_loss: 0.4926 - val_acc: 0.8462
Epoch 34/50
40000/40000 - 11s - loss: 0.2432 - acc: 0.9157 - val_loss: 0.5137 - val_acc: 0.8429
Epoch 35/50
40000/40000 - 11s - loss: 0.2414 - acc: 0.9144 - val_loss: 0.5184 - val_acc: 0.8430
Epoch 36/50
40000/40000 - 11s - loss: 0.2319 - acc: 0.9173 - val_loss: 0.5304 - val_acc: 0.8451
Epoch 37/50
40000/40000 - 11s - loss: 0.2241 - acc: 0.9207 - val_loss: 0.5480 - val_acc: 0.8421
Epoch 38/50
40000/40000 - 11s - loss: 0.2213 - acc: 0.9213 - val_loss: 0.5259 - val_acc: 0.8492
Epoch 39/50
40000/40000 - 11s - loss: 0.2194 - acc: 0.9230 - val_loss: 0.5425 - val_acc: 0.8415
Epoch 40/50
40000/40000 - 11s - loss: 0.2139 - acc: 0.9245 - val_loss: 0.5644 - val_acc: 0.8389
Epoch 41/50
40000/40000 - 11s - loss: 0.2119 - acc: 0.9262 - val_loss: 0.5224 - val_acc:

```

0.8475
Epoch 42/50
40000/40000 - 11s - loss: 0.2064 - acc: 0.9263 - val_loss: 0.5293 - val_acc:
0.8477
Epoch 43/50
40000/40000 - 11s - loss: 0.1942 - acc: 0.9309 - val_loss: 0.5854 - val_acc:
0.8330
Epoch 44/50
40000/40000 - 11s - loss: 0.1969 - acc: 0.9313 - val_loss: 0.5293 - val_acc:
0.8499
Epoch 45/50
40000/40000 - 11s - loss: 0.1918 - acc: 0.9326 - val_loss: 0.5236 - val_acc:
0.8505
Epoch 46/50
40000/40000 - 11s - loss: 0.1869 - acc: 0.9331 - val_loss: 0.5280 - val_acc:
0.8508
Epoch 47/50
40000/40000 - 11s - loss: 0.1831 - acc: 0.9360 - val_loss: 0.5356 - val_acc:
0.8486
Epoch 48/50
40000/40000 - 11s - loss: 0.1847 - acc: 0.9352 - val_loss: 0.5336 - val_acc:
0.8486
Epoch 49/50
40000/40000 - 11s - loss: 0.1772 - acc: 0.9367 - val_loss: 0.5371 - val_acc:
0.8508
Epoch 50/50
40000/40000 - 11s - loss: 0.1842 - acc: 0.9358 - val_loss: 0.5303 - val_acc:
0.8479

```

1.9.3 Adagrad

```

[33]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',

```

```

        kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.Adagrad(learning_rate=0.01)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history3 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model

```



```
model.save('model_6_part3.h5', overwrite=True)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/keras/optimizer_v2/adagrad.py:107: calling Constant.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

Train on 40000 samples, validate on 10000 samples

Epoch 1/50

40000/40000 - 14s - loss: 1.7443 - acc: 0.4086 - val_loss: 1.1958 - val_acc: 0.5661

Epoch 2/50

40000/40000 - 11s - loss: 1.2376 - acc: 0.5635 - val_loss: 1.0010 - val_acc: 0.6428

Epoch 3/50

40000/40000 - 11s - loss: 1.0434 - acc: 0.6277 - val_loss: 0.9030 - val_acc: 0.6806

Epoch 4/50

40000/40000 - 11s - loss: 0.9333 - acc: 0.6696 - val_loss: 0.8655 - val_acc: 0.6870

Epoch 5/50

40000/40000 - 11s - loss: 0.8611 - acc: 0.6963 - val_loss: 0.7300 - val_acc: 0.7430

Epoch 6/50

40000/40000 - 11s - loss: 0.7979 - acc: 0.7198 - val_loss: 0.7097 - val_acc: 0.7463

Epoch 7/50

40000/40000 - 11s - loss: 0.7545 - acc: 0.7372 - val_loss: 0.6592 - val_acc: 0.7662

Epoch 8/50

40000/40000 - 11s - loss: 0.7194 - acc: 0.7486 - val_loss: 0.6989 - val_acc: 0.7541

Epoch 9/50

40000/40000 - 11s - loss: 0.6848 - acc: 0.7596 - val_loss: 0.6479 - val_acc: 0.7731

Epoch 10/50

40000/40000 - 11s - loss: 0.6593 - acc: 0.7691 - val_loss: 0.6236 - val_acc: 0.7756

Epoch 11/50

40000/40000 - 11s - loss: 0.6315 - acc: 0.7789 - val_loss: 0.5948 - val_acc: 0.7866

Epoch 12/50

40000/40000 - 11s - loss: 0.6144 - acc: 0.7870 - val_loss: 0.5938 - val_acc: 0.7936

Epoch 13/50

40000/40000 - 11s - loss: 0.5892 - acc: 0.7926 - val_loss: 0.5849 - val_acc: 0.7922
Epoch 14/50
40000/40000 - 11s - loss: 0.5647 - acc: 0.8026 - val_loss: 0.5564 - val_acc: 0.8021
Epoch 15/50
40000/40000 - 11s - loss: 0.5514 - acc: 0.8062 - val_loss: 0.5543 - val_acc: 0.8079
Epoch 16/50
40000/40000 - 11s - loss: 0.5351 - acc: 0.8125 - val_loss: 0.5460 - val_acc: 0.8093
Epoch 17/50
40000/40000 - 11s - loss: 0.5205 - acc: 0.8192 - val_loss: 0.5562 - val_acc: 0.8082
Epoch 18/50
40000/40000 - 11s - loss: 0.5066 - acc: 0.8255 - val_loss: 0.5189 - val_acc: 0.8208
Epoch 19/50
40000/40000 - 11s - loss: 0.4895 - acc: 0.8285 - val_loss: 0.5255 - val_acc: 0.8189
Epoch 20/50
40000/40000 - 11s - loss: 0.4760 - acc: 0.8337 - val_loss: 0.5282 - val_acc: 0.8207
Epoch 21/50
40000/40000 - 11s - loss: 0.4620 - acc: 0.8392 - val_loss: 0.5274 - val_acc: 0.8187
Epoch 22/50
40000/40000 - 11s - loss: 0.4535 - acc: 0.8402 - val_loss: 0.5203 - val_acc: 0.8223
Epoch 23/50
40000/40000 - 11s - loss: 0.4370 - acc: 0.8471 - val_loss: 0.5064 - val_acc: 0.8267
Epoch 24/50
40000/40000 - 11s - loss: 0.4344 - acc: 0.8474 - val_loss: 0.5283 - val_acc: 0.8212
Epoch 25/50
40000/40000 - 11s - loss: 0.4255 - acc: 0.8499 - val_loss: 0.5124 - val_acc: 0.8254
Epoch 26/50
40000/40000 - 11s - loss: 0.4102 - acc: 0.8574 - val_loss: 0.5051 - val_acc: 0.8284
Epoch 27/50
40000/40000 - 11s - loss: 0.4009 - acc: 0.8582 - val_loss: 0.5119 - val_acc: 0.8291
Epoch 28/50
40000/40000 - 11s - loss: 0.3886 - acc: 0.8637 - val_loss: 0.5042 - val_acc: 0.8296
Epoch 29/50

40000/40000 - 11s - loss: 0.3837 - acc: 0.8654 - val_loss: 0.4996 - val_acc: 0.8344
Epoch 30/50
40000/40000 - 11s - loss: 0.3725 - acc: 0.8672 - val_loss: 0.5034 - val_acc: 0.8320
Epoch 31/50
40000/40000 - 11s - loss: 0.3674 - acc: 0.8701 - val_loss: 0.4858 - val_acc: 0.8377
Epoch 32/50
40000/40000 - 11s - loss: 0.3519 - acc: 0.8765 - val_loss: 0.4984 - val_acc: 0.8354
Epoch 33/50
40000/40000 - 11s - loss: 0.3530 - acc: 0.8757 - val_loss: 0.5067 - val_acc: 0.8316
Epoch 34/50
40000/40000 - 11s - loss: 0.3425 - acc: 0.8793 - val_loss: 0.4968 - val_acc: 0.8351
Epoch 35/50
40000/40000 - 11s - loss: 0.3347 - acc: 0.8813 - val_loss: 0.5172 - val_acc: 0.8325
Epoch 36/50
40000/40000 - 11s - loss: 0.3324 - acc: 0.8838 - val_loss: 0.4867 - val_acc: 0.8433
Epoch 37/50
40000/40000 - 11s - loss: 0.3244 - acc: 0.8859 - val_loss: 0.4922 - val_acc: 0.8407
Epoch 38/50
40000/40000 - 11s - loss: 0.3117 - acc: 0.8906 - val_loss: 0.4969 - val_acc: 0.8421
Epoch 39/50
40000/40000 - 11s - loss: 0.3096 - acc: 0.8890 - val_loss: 0.4969 - val_acc: 0.8432
Epoch 40/50
40000/40000 - 11s - loss: 0.3002 - acc: 0.8946 - val_loss: 0.4989 - val_acc: 0.8398
Epoch 41/50
40000/40000 - 11s - loss: 0.2980 - acc: 0.8958 - val_loss: 0.4880 - val_acc: 0.8432
Epoch 42/50
40000/40000 - 11s - loss: 0.2921 - acc: 0.8976 - val_loss: 0.4988 - val_acc: 0.8422
Epoch 43/50
40000/40000 - 11s - loss: 0.2865 - acc: 0.9004 - val_loss: 0.5134 - val_acc: 0.8356
Epoch 44/50
40000/40000 - 11s - loss: 0.2809 - acc: 0.9018 - val_loss: 0.4947 - val_acc: 0.8444
Epoch 45/50

```

40000/40000 - 11s - loss: 0.2729 - acc: 0.9044 - val_loss: 0.5020 - val_acc:
0.8424
Epoch 46/50
40000/40000 - 11s - loss: 0.2689 - acc: 0.9068 - val_loss: 0.5023 - val_acc:
0.8399
Epoch 47/50
40000/40000 - 11s - loss: 0.2634 - acc: 0.9075 - val_loss: 0.4968 - val_acc:
0.8431
Epoch 48/50
40000/40000 - 11s - loss: 0.2592 - acc: 0.9096 - val_loss: 0.4944 - val_acc:
0.8433
Epoch 49/50
40000/40000 - 11s - loss: 0.2557 - acc: 0.9089 - val_loss: 0.5066 - val_acc:
0.8426
Epoch 50/50
40000/40000 - 11s - loss: 0.2571 - acc: 0.9094 - val_loss: 0.5079 - val_acc:
0.8434

```

1.9.4 Adadelta

A robust extension of Adagrad that adapts learning rates based on a moving window of gradient updates (source <https://keras.io/optimizers/>)

```

[34]: model = Sequential()
      model.add(Conv2D(32, (3, 3), padding='same',
                      kernel_initializer='random_uniform',
                      input_shape=shape))
      model.add(Activation('relu'))
      model.add(BatchNormalization())
      model.add(Conv2D(32, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(BatchNormalization())
      model.add(MaxPooling2D(pool_size=(2, 2)))
      model.add(Dropout(0.15))

      model.add(Conv2D(64, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(BatchNormalization())
      model.add(Conv2D(64, (3, 3),
                      padding='same',
                      kernel_initializer='random_uniform'))
      model.add(Activation('relu'))
      model.add(BatchNormalization())
      model.add(MaxPooling2D(pool_size=(2, 2)))

```

```

model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Flatten())
model.add(Dense(256))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.45))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.Adadelta(learning_rate=1.0, rho=0.95)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history4 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_6_part4.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/50

40000/40000 - 14s - loss: 1.6527 - acc: 0.4318 - val_loss: 1.2876 - val_acc: 0.5391

Epoch 2/50

40000/40000 - 11s - loss: 1.0363 - acc: 0.6357 - val_loss: 1.2510 - val_acc: 0.6011

Epoch 3/50
40000/40000 - 11s - loss: 0.8458 - acc: 0.7068 - val_loss: 0.7598 - val_acc: 0.7344
Epoch 4/50
40000/40000 - 11s - loss: 0.7398 - acc: 0.7442 - val_loss: 0.6771 - val_acc: 0.7661
Epoch 5/50
40000/40000 - 11s - loss: 0.6627 - acc: 0.7688 - val_loss: 0.7105 - val_acc: 0.7532
Epoch 6/50
40000/40000 - 11s - loss: 0.6178 - acc: 0.7850 - val_loss: 0.6000 - val_acc: 0.7932
Epoch 7/50
40000/40000 - 11s - loss: 0.5658 - acc: 0.8049 - val_loss: 0.5570 - val_acc: 0.8109
Epoch 8/50
40000/40000 - 11s - loss: 0.5270 - acc: 0.8176 - val_loss: 0.5395 - val_acc: 0.8187
Epoch 9/50
40000/40000 - 11s - loss: 0.4941 - acc: 0.8307 - val_loss: 0.5526 - val_acc: 0.8123
Epoch 10/50
40000/40000 - 11s - loss: 0.4726 - acc: 0.8379 - val_loss: 0.5618 - val_acc: 0.8110
Epoch 11/50
40000/40000 - 11s - loss: 0.4443 - acc: 0.8475 - val_loss: 0.5709 - val_acc: 0.8114
Epoch 12/50
40000/40000 - 11s - loss: 0.4203 - acc: 0.8557 - val_loss: 0.5107 - val_acc: 0.8330
Epoch 13/50
40000/40000 - 11s - loss: 0.4035 - acc: 0.8623 - val_loss: 0.5200 - val_acc: 0.8287
Epoch 14/50
40000/40000 - 11s - loss: 0.3859 - acc: 0.8670 - val_loss: 0.5502 - val_acc: 0.8262
Epoch 15/50
40000/40000 - 11s - loss: 0.3711 - acc: 0.8714 - val_loss: 0.5321 - val_acc: 0.8315
Epoch 16/50
40000/40000 - 11s - loss: 0.3503 - acc: 0.8789 - val_loss: 0.4910 - val_acc: 0.8417
Epoch 17/50
40000/40000 - 11s - loss: 0.3402 - acc: 0.8831 - val_loss: 0.5171 - val_acc: 0.8364
Epoch 18/50
40000/40000 - 11s - loss: 0.3265 - acc: 0.8860 - val_loss: 0.5020 - val_acc: 0.8373

Epoch 19/50
40000/40000 - 11s - loss: 0.3121 - acc: 0.8915 - val_loss: 0.5363 - val_acc: 0.8391
Epoch 20/50
40000/40000 - 11s - loss: 0.2993 - acc: 0.8956 - val_loss: 0.5646 - val_acc: 0.8231
Epoch 21/50
40000/40000 - 11s - loss: 0.2847 - acc: 0.9020 - val_loss: 0.4825 - val_acc: 0.8510
Epoch 22/50
40000/40000 - 11s - loss: 0.2776 - acc: 0.9036 - val_loss: 0.5395 - val_acc: 0.8403
Epoch 23/50
40000/40000 - 11s - loss: 0.2743 - acc: 0.9063 - val_loss: 0.4968 - val_acc: 0.8508
Epoch 24/50
40000/40000 - 11s - loss: 0.2637 - acc: 0.9086 - val_loss: 0.4972 - val_acc: 0.8458
Epoch 25/50
40000/40000 - 12s - loss: 0.2583 - acc: 0.9098 - val_loss: 0.5082 - val_acc: 0.8468
Epoch 26/50
40000/40000 - 11s - loss: 0.2534 - acc: 0.9117 - val_loss: 0.5320 - val_acc: 0.8427
Epoch 27/50
40000/40000 - 11s - loss: 0.2458 - acc: 0.9143 - val_loss: 0.5404 - val_acc: 0.8434
Epoch 28/50
40000/40000 - 11s - loss: 0.2371 - acc: 0.9156 - val_loss: 0.5814 - val_acc: 0.8276
Epoch 29/50
40000/40000 - 11s - loss: 0.2320 - acc: 0.9208 - val_loss: 0.5013 - val_acc: 0.8544
Epoch 30/50
40000/40000 - 11s - loss: 0.2275 - acc: 0.9227 - val_loss: 0.5433 - val_acc: 0.8436
Epoch 31/50
40000/40000 - 11s - loss: 0.2180 - acc: 0.9232 - val_loss: 0.5044 - val_acc: 0.8472
Epoch 32/50
40000/40000 - 11s - loss: 0.2220 - acc: 0.9236 - val_loss: 0.5363 - val_acc: 0.8463
Epoch 33/50
40000/40000 - 11s - loss: 0.2143 - acc: 0.9259 - val_loss: 0.5084 - val_acc: 0.8565
Epoch 34/50
40000/40000 - 11s - loss: 0.2097 - acc: 0.9272 - val_loss: 0.5568 - val_acc: 0.8501

Epoch 35/50
40000/40000 - 11s - loss: 0.2025 - acc: 0.9285 - val_loss: 0.5613 - val_acc: 0.8474
Epoch 36/50
40000/40000 - 11s - loss: 0.1968 - acc: 0.9324 - val_loss: 0.5926 - val_acc: 0.8372
Epoch 37/50
40000/40000 - 11s - loss: 0.2007 - acc: 0.9306 - val_loss: 0.5479 - val_acc: 0.8485
Epoch 38/50
40000/40000 - 11s - loss: 0.1877 - acc: 0.9364 - val_loss: 0.5358 - val_acc: 0.8560
Epoch 39/50
40000/40000 - 11s - loss: 0.1918 - acc: 0.9350 - val_loss: 0.5357 - val_acc: 0.8490
Epoch 40/50
40000/40000 - 11s - loss: 0.1891 - acc: 0.9354 - val_loss: 0.5832 - val_acc: 0.8414
Epoch 41/50
40000/40000 - 11s - loss: 0.1877 - acc: 0.9356 - val_loss: 0.5358 - val_acc: 0.8520
Epoch 42/50
40000/40000 - 11s - loss: 0.1756 - acc: 0.9390 - val_loss: 0.5488 - val_acc: 0.8492
Epoch 43/50
40000/40000 - 11s - loss: 0.1803 - acc: 0.9386 - val_loss: 0.5241 - val_acc: 0.8595
Epoch 44/50
40000/40000 - 11s - loss: 0.1733 - acc: 0.9403 - val_loss: 0.5071 - val_acc: 0.8530
Epoch 45/50
40000/40000 - 11s - loss: 0.1651 - acc: 0.9423 - val_loss: 0.5281 - val_acc: 0.8617
Epoch 46/50
40000/40000 - 11s - loss: 0.1699 - acc: 0.9407 - val_loss: 0.5472 - val_acc: 0.8531
Epoch 47/50
40000/40000 - 11s - loss: 0.1614 - acc: 0.9442 - val_loss: 0.5655 - val_acc: 0.8527
Epoch 48/50
40000/40000 - 11s - loss: 0.1637 - acc: 0.9434 - val_loss: 0.5456 - val_acc: 0.8527
Epoch 49/50
40000/40000 - 11s - loss: 0.1667 - acc: 0.9421 - val_loss: 0.5326 - val_acc: 0.8531
Epoch 50/50
40000/40000 - 11s - loss: 0.1609 - acc: 0.9449 - val_loss: 0.5690 - val_acc: 0.8490

1.9.5 Validation

```
[36]: print("Optimizers/n")


print("RMSprop")
plot_graphs(history1)


print("SGD")
plot_graphs(history2)

print("Adagrad")
plot_graphs(history3)


print("Adadelat")
plot_graphs(history4)
```


Optimizers/n
RMSprop







SGD







Adagrad



Adadelta



1.9.6 Validation Results

From the graphs as well as the output per model, we can see that all optimizers tend to start overfitting after around 25 epochs. To counteract this, we need to increase the regularization (dropout layer) or we can limit the epochs to 25. Since the accuracy doesn't seem to increase much after 25, we will keep it around 25.

At around epoch 25, all the optimizers seem to perform very similarly, with small differences with increase in accuracy for slightly more overfitted models. To create our final models, we will stick to 30 for number of epochs and RMSprop and Adadelta for our optimizers (best fit model and most accurate model) and slightly increase the regularization.

1.10 Model #7 - Final touches

For our final models, we will increase the number of layers to see if that increases our accuracy and reduces the training/validation loss. Along with the increased number of layers, based on model #6, we will increase the number of epochs slightly to 30 along with slightly higher regularization (dropout increases from 0.25 -> 0.55)

```
[0]: epochs = 30
```

```
[70]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```

model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Conv2D(256, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(256, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.45))

model.add(Conv2D(512, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(512, (3, 3),
                 padding='same',
                 kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.55))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.55))
model.add(Dense(num_classes))

```

```

model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])

#Train model
history1 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_7_part1.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

Epoch 1/30

40000/40000 - 36s - loss: 2.8870 - acc: 0.1779 - val_loss: 1.9347 - val_acc: 0.2806

Epoch 2/30

40000/40000 - 24s - loss: 2.0701 - acc: 0.3055 - val_loss: 1.5718 - val_acc: 0.4171

Epoch 3/30

40000/40000 - 25s - loss: 1.6754 - acc: 0.4223 - val_loss: 1.3655 - val_acc: 0.5190

Epoch 4/30

40000/40000 - 25s - loss: 1.4175 - acc: 0.5077 - val_loss: 1.3039 - val_acc: 0.5410

Epoch 5/30

40000/40000 - 25s - loss: 1.2404 - acc: 0.5747 - val_loss: 1.0291 - val_acc: 0.6402

Epoch 6/30

40000/40000 - 25s - loss: 1.1190 - acc: 0.6182 - val_loss: 0.9412 - val_acc: 0.6748

Epoch 7/30

40000/40000 - 25s - loss: 1.0191 - acc: 0.6506 - val_loss: 0.8903 - val_acc: 0.6860

Epoch 8/30

40000/40000 - 25s - loss: 0.9466 - acc: 0.6790 - val_loss: 0.7647 - val_acc: 0.7351

Epoch 9/30

40000/40000 - 24s - loss: 0.8887 - acc: 0.7010 - val_loss: 0.7762 - val_acc: 0.7299

Epoch 10/30
40000/40000 - 25s - loss: 0.8385 - acc: 0.7187 - val_loss: 0.7193 - val_acc: 0.7529
Epoch 11/30
40000/40000 - 25s - loss: 0.7910 - acc: 0.7353 - val_loss: 0.7048 - val_acc: 0.7597
Epoch 12/30
40000/40000 - 25s - loss: 0.7547 - acc: 0.7483 - val_loss: 0.6544 - val_acc: 0.7776
Epoch 13/30
40000/40000 - 25s - loss: 0.7214 - acc: 0.7585 - val_loss: 0.7536 - val_acc: 0.7537
Epoch 14/30
40000/40000 - 24s - loss: 0.6889 - acc: 0.7692 - val_loss: 0.6857 - val_acc: 0.7720
Epoch 15/30
40000/40000 - 24s - loss: 0.6523 - acc: 0.7836 - val_loss: 0.5984 - val_acc: 0.7982
Epoch 16/30
40000/40000 - 25s - loss: 0.6256 - acc: 0.7937 - val_loss: 0.5703 - val_acc: 0.8098
Epoch 17/30
40000/40000 - 25s - loss: 0.6172 - acc: 0.7968 - val_loss: 0.6295 - val_acc: 0.7896
Epoch 18/30
40000/40000 - 25s - loss: 0.5907 - acc: 0.8069 - val_loss: 0.5685 - val_acc: 0.8105
Epoch 19/30
40000/40000 - 25s - loss: 0.5732 - acc: 0.8111 - val_loss: 0.5549 - val_acc: 0.8194
Epoch 20/30
40000/40000 - 25s - loss: 0.5493 - acc: 0.8195 - val_loss: 0.5348 - val_acc: 0.8194
Epoch 21/30
40000/40000 - 25s - loss: 0.5338 - acc: 0.8266 - val_loss: 0.5462 - val_acc: 0.8202
Epoch 22/30
40000/40000 - 25s - loss: 0.5148 - acc: 0.8315 - val_loss: 0.5565 - val_acc: 0.8171
Epoch 23/30
40000/40000 - 25s - loss: 0.4995 - acc: 0.8362 - val_loss: 0.5320 - val_acc: 0.8241
Epoch 24/30
40000/40000 - 25s - loss: 0.4877 - acc: 0.8418 - val_loss: 0.5564 - val_acc: 0.8207
Epoch 25/30
40000/40000 - 25s - loss: 0.4763 - acc: 0.8466 - val_loss: 0.4911 - val_acc: 0.8417

Epoch 26/30
 40000/40000 - 25s - loss: 0.4530 - acc: 0.8534 - val_loss: 0.6040 - val_acc: 0.8073
 Epoch 27/30
 40000/40000 - 25s - loss: 0.4447 - acc: 0.8556 - val_loss: 0.5570 - val_acc: 0.8210
 Epoch 28/30
 40000/40000 - 25s - loss: 0.4236 - acc: 0.8624 - val_loss: 0.5740 - val_acc: 0.8232
 Epoch 29/30
 40000/40000 - 24s - loss: 0.4209 - acc: 0.8632 - val_loss: 0.5200 - val_acc: 0.8390
 Epoch 30/30
 40000/40000 - 25s - loss: 0.4073 - acc: 0.8682 - val_loss: 0.4936 - val_acc: 0.8437

```
[73]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                kernel_initializer='random_uniform',
                input_shape=shape))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.15))

model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
```

```

model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.35))

model.add(Conv2D(256, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(256, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.45))

model.add(Conv2D(512, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Conv2D(512, (3, 3),
                padding='same',
                kernel_initializer='random_uniform'))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.55))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.55))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

# Compile the model
opt = keras.optimizers.Adadelta(learning_rate=1.0, rho=0.95)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,

```

```

        metrics=['accuracy'])

#Train model
history2 = model.fit(x_train,
                    to_categorical(y_train, num_classes),
                    epochs=epochs,
                    verbose=2,
                    validation_split=0.2,
                    shuffle=True)

#Save model
model.save('model_7_part2.h5', overwrite=True)

```

Train on 40000 samples, validate on 10000 samples

```

Epoch 1/30
40000/40000 - 36s - loss: 2.0485 - acc: 0.3242 - val_loss: 1.6751 - val_acc:
0.4436
Epoch 2/30
40000/40000 - 23s - loss: 1.2756 - acc: 0.5503 - val_loss: 1.0582 - val_acc:
0.6302
Epoch 3/30
40000/40000 - 23s - loss: 1.0445 - acc: 0.6426 - val_loss: 0.9104 - val_acc:
0.6892
Epoch 4/30
40000/40000 - 23s - loss: 0.8943 - acc: 0.6961 - val_loss: 0.7240 - val_acc:
0.7513
Epoch 5/30
40000/40000 - 23s - loss: 0.7978 - acc: 0.7339 - val_loss: 0.6624 - val_acc:
0.7795
Epoch 6/30
40000/40000 - 23s - loss: 0.7175 - acc: 0.7616 - val_loss: 0.6401 - val_acc:
0.7817
Epoch 7/30
40000/40000 - 23s - loss: 0.6604 - acc: 0.7804 - val_loss: 0.5858 - val_acc:
0.8040
Epoch 8/30
40000/40000 - 23s - loss: 0.6089 - acc: 0.7995 - val_loss: 0.6338 - val_acc:
0.7917
Epoch 9/30
40000/40000 - 23s - loss: 0.5638 - acc: 0.8143 - val_loss: 0.5560 - val_acc:
0.8157
Epoch 10/30
40000/40000 - 23s - loss: 0.5292 - acc: 0.8255 - val_loss: 0.5361 - val_acc:
0.8158
Epoch 11/30
40000/40000 - 23s - loss: 0.4995 - acc: 0.8352 - val_loss: 0.5433 - val_acc:
0.8184

```

Epoch 12/30
40000/40000 - 23s - loss: 0.4661 - acc: 0.8478 - val_loss: 0.5173 - val_acc: 0.8312

Epoch 13/30
40000/40000 - 23s - loss: 0.4432 - acc: 0.8544 - val_loss: 0.5869 - val_acc: 0.8107

Epoch 14/30
40000/40000 - 23s - loss: 0.4177 - acc: 0.8634 - val_loss: 0.5087 - val_acc: 0.8316

Epoch 15/30
40000/40000 - 23s - loss: 0.3956 - acc: 0.8687 - val_loss: 0.6058 - val_acc: 0.8141

Epoch 16/30
40000/40000 - 23s - loss: 0.3772 - acc: 0.8745 - val_loss: 0.5051 - val_acc: 0.8367

Epoch 17/30
40000/40000 - 23s - loss: 0.3570 - acc: 0.8813 - val_loss: 0.4805 - val_acc: 0.8471

Epoch 18/30
40000/40000 - 23s - loss: 0.3386 - acc: 0.8882 - val_loss: 0.5333 - val_acc: 0.8385

Epoch 19/30
40000/40000 - 23s - loss: 0.3305 - acc: 0.8903 - val_loss: 0.4965 - val_acc: 0.8454

Epoch 20/30
40000/40000 - 22s - loss: 0.3110 - acc: 0.8972 - val_loss: 0.4838 - val_acc: 0.8525

Epoch 21/30
40000/40000 - 23s - loss: 0.2988 - acc: 0.9022 - val_loss: 0.5181 - val_acc: 0.8424

Epoch 22/30
40000/40000 - 23s - loss: 0.2868 - acc: 0.9075 - val_loss: 0.5428 - val_acc: 0.8396

Epoch 23/30
40000/40000 - 23s - loss: 0.2795 - acc: 0.9075 - val_loss: 0.5098 - val_acc: 0.8466

Epoch 24/30
40000/40000 - 23s - loss: 0.2619 - acc: 0.9128 - val_loss: 0.5514 - val_acc: 0.8351

Epoch 25/30
40000/40000 - 23s - loss: 0.2532 - acc: 0.9188 - val_loss: 0.5214 - val_acc: 0.8516

Epoch 26/30
40000/40000 - 23s - loss: 0.2503 - acc: 0.9183 - val_loss: 0.4986 - val_acc: 0.8493


Epoch 27/30
40000/40000 - 22s - loss: 0.2406 - acc: 0.9194 - val_loss: 0.5347 - val_acc: 0.8470


Epoch 28/30
40000/40000 - 22s - loss: 0.2307 - acc: 0.9241 - val_loss: 0.5404 - val_acc: 0.8401
Epoch 29/30
40000/40000 - 23s - loss: 0.2256 - acc: 0.9261 - val_loss: 0.5513 - val_acc: 0.8466
Epoch 30/30
40000/40000 - 23s - loss: 0.2129 - acc: 0.9307 - val_loss: 0.5880 - val_acc: 0.8449

1.10.1 Validation


```
[74]: print("RMSprop")  
      plot_graphs(history1)  
  
      print("Adadelata")  
      plot_graphs(history2)
```


RMSprop





Adadelta







1.11 Final Model (Choosing the best one)

The final model we will choose is model #7 part 1. This model uses the optimizer RMSprop, uses an increasing dropout of 0.25 -> 0.55 regularizer, Batch Normalization after every layer, 30 epochs with batch size 32, and 10 cnn layers followed by a flattening layer.

Net architecture is printed below along with training/validation accuracy and loss graph.

```
[77]: plot_graphs(history1)
      model = load_model("model_7_part1.h5")
      print(model.summary())
```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
conv2d_204 (Conv2D)	(None, 32, 32, 32)	896
activation_254 (Activation)	(None, 32, 32, 32)	0
batch_normalization_220 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_205 (Conv2D)	(None, 32, 32, 32)	9248
activation_255 (Activation)	(None, 32, 32, 32)	0
batch_normalization_221 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_101 (Max Pooling)	(None, 16, 16, 32)	0
dropout_122 (Dropout)	(None, 16, 16, 32)	0
conv2d_206 (Conv2D)	(None, 16, 16, 64)	18496
activation_256 (Activation)	(None, 16, 16, 64)	0
batch_normalization_222 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_207 (Conv2D)	(None, 16, 16, 64)	36928
activation_257 (Activation)	(None, 16, 16, 64)	0
batch_normalization_223 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_102 (Max Pooling)	(None, 8, 8, 64)	0
dropout_123 (Dropout)	(None, 8, 8, 64)	0
conv2d_208 (Conv2D)	(None, 8, 8, 128)	73856
activation_258 (Activation)	(None, 8, 8, 128)	0
batch_normalization_224 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_209 (Conv2D)	(None, 8, 8, 128)	147584
activation_259 (Activation)	(None, 8, 8, 128)	0
batch_normalization_225 (Batch Normalization)	(None, 8, 8, 128)	512

max_pooling2d_103 (MaxPoolin	(None, 4, 4, 128)	0
dropout_124 (Dropout)	(None, 4, 4, 128)	0
conv2d_210 (Conv2D)	(None, 4, 4, 256)	295168
activation_260 (Activation)	(None, 4, 4, 256)	0
batch_normalization_226 (Bat	(None, 4, 4, 256)	1024
conv2d_211 (Conv2D)	(None, 4, 4, 256)	590080
activation_261 (Activation)	(None, 4, 4, 256)	0
batch_normalization_227 (Bat	(None, 4, 4, 256)	1024
max_pooling2d_104 (MaxPoolin	(None, 2, 2, 256)	0
dropout_125 (Dropout)	(None, 2, 2, 256)	0
conv2d_212 (Conv2D)	(None, 2, 2, 512)	1180160
activation_262 (Activation)	(None, 2, 2, 512)	0
batch_normalization_228 (Bat	(None, 2, 2, 512)	2048
conv2d_213 (Conv2D)	(None, 2, 2, 512)	2359808
activation_263 (Activation)	(None, 2, 2, 512)	0
batch_normalization_229 (Bat	(None, 2, 2, 512)	2048
max_pooling2d_105 (MaxPoolin	(None, 1, 1, 512)	0
dropout_126 (Dropout)	(None, 1, 1, 512)	0
flatten_25 (Flatten)	(None, 512)	0
dense_50 (Dense)	(None, 128)	65664
activation_264 (Activation)	(None, 128)	0
batch_normalization_230 (Bat	(None, 128)	512
dropout_127 (Dropout)	(None, 128)	0
dense_51 (Dense)	(None, 10)	1290

```

activation_265 (Activation) (None, 10) 0
=====
Total params: 4,787,626
Trainable params: 4,783,402
Non-trainable params: 4,224
-----
None

```

1.12 Testing Against Final Model

We will finally test our best model using the test data that we saved in the beginning. The following is the classification report and confusion matrix (created using helper functions in the second defined in the second section - useful helper functions).

```
[80]: get_metrics("model_7_part1.h5")
```

```

10000/10000 [=====] - 6s 626us/sample - loss: 0.5163 -
acc: 0.8375

```

```
[0.5162984648227692, 0.8375]
```

	precision	recall	f1-score	support
airplane	0.83	0.86	0.84	1000
automobile	0.92	0.93	0.92	1000
bird	0.82	0.72	0.77	1000
cat	0.71	0.64	0.68	1000
deer	0.82	0.85	0.84	1000
dog	0.75	0.77	0.76	1000
frog	0.84	0.90	0.87	1000
horse	0.88	0.88	0.88	1000
ship	0.91	0.89	0.90	1000
truck	0.87	0.93	0.90	1000
accuracy			0.84	10000
macro avg	0.84	0.84	0.84	10000
weighted avg	0.84	0.84	0.84	10000

	predicted - airplane	predicted - automobile	predicted - bird \
airplane	862	13	16
automobile	3	926	3
bird	65	0	723
cat	22	1	41
deer	10	1	34
dog	9	3	27
frog	5	2	18
horse	11	2	16
ship	44	22	2
truck	10	36	1

	predicted - cat	predicted - deer	predicted - dog \
airplane	8	6	2
automobile	1	2	1
bird	39	54	44
cat	643	39	150
deer	30	851	14
dog	113	30	765
frog	32	19	11
horse	25	28	29
ship	8	3	0
truck	6	0	0

	predicted - frog	predicted - horse	predicted - ship \
airplane	9	6	42
automobile	3	0	6
bird	55	10	5
cat	61	26	11
deer	23	33	4
dog	13	35	1
frog	902	7	4
horse	2	880	1
ship	3	0	892
truck	4	3	9

	predicted - truck
airplane	36
automobile	55
bird	5
cat	6
deer	0
dog	4
frog	0
horse	6
ship	26
truck	931

2 Conclusion/Results

In this project, we started by processing our data set and dividing our data into our training set and our test set. With our training set, we divided this into 80% training and 20% validation for each model (using the `validation_split=0.2` variable when training the training data).

In our first model, we created a 6 layer cnn along with a flatten layer. We used an arbitrary optimizer (RMSprop) to begin with. As seen by the training/validation loss graph, our model was highly overfitting the data. Thus, regularization was needed.

With our second, third, and fourth models, we tried various parameters to regularize our data. In our second model, we tried using different weight decays to regularize (0, 0.00005, 0.0005, and 0.05). In our third model, we tried using different dropout layers (0.25, 0.5, and an increasing 0.15->0.45). In our fourth model, we used a combination of weight decay and dropout. With all these regularization techniques, it was determined that either the increasing dropout layer or the increasing dropout layer plus weight decay of 0.0005 best prevented overfitting while keeping accuracy up.

In our fifth model, we added normalization using Batch Normalization. Batch Normalization was added to each previous hidden layer at each batch and worked to reduce some overfitting and improve accuracy overall.

In our sixth model, we tried using different optimizers and epoch size 50. In the end, epochs of size 50 started overfitting the data and the good mixture of accuracy and preventing overfitting seemed to be around 30 epochs. The optimizers seemed to be relatively similar in performance, with RMSprop having the least difference between the training and validation loss and AdaDelta having the best overall accuracy on the validation data.

In our seventh and final model sets, we tried adjusting the epoch number to 35 as well as increasing the 6 layer cnn to 10 layers (plus the flatten layer).

Overall, the seventh model's first model seemed to have the best mixture of accuracy and minimizing the training and validation model loss. Thus, we chose this as our final model and ran this against our test set that we saved in the beginning.

2.1 Results

As stated in the Final Model section, the final model we chose and ran our test against is model #7 part 1. This model uses the optimizer RMSprop, uses an increasing dropout of 0.25 -> 0.55 for regularization, Batch Normalization after every layer, 30 epochs with batch size 32, and 10 cnn layers followed by a flattening layer.

The net architecture is as following:

```
[81]: model = load_model("model_7_part1.h5")
      print(model.summary())
```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
conv2d_204 (Conv2D)	(None, 32, 32, 32)	896
activation_254 (Activation)	(None, 32, 32, 32)	0
batch_normalization_220 (Bat	(None, 32, 32, 32)	128

conv2d_205 (Conv2D)	(None, 32, 32, 32)	9248
activation_255 (Activation)	(None, 32, 32, 32)	0
batch_normalization_221 (Bat	(None, 32, 32, 32)	128
max_pooling2d_101 (MaxPoolin	(None, 16, 16, 32)	0
dropout_122 (Dropout)	(None, 16, 16, 32)	0
conv2d_206 (Conv2D)	(None, 16, 16, 64)	18496
activation_256 (Activation)	(None, 16, 16, 64)	0
batch_normalization_222 (Bat	(None, 16, 16, 64)	256
conv2d_207 (Conv2D)	(None, 16, 16, 64)	36928
activation_257 (Activation)	(None, 16, 16, 64)	0
batch_normalization_223 (Bat	(None, 16, 16, 64)	256
max_pooling2d_102 (MaxPoolin	(None, 8, 8, 64)	0
dropout_123 (Dropout)	(None, 8, 8, 64)	0
conv2d_208 (Conv2D)	(None, 8, 8, 128)	73856
activation_258 (Activation)	(None, 8, 8, 128)	0
batch_normalization_224 (Bat	(None, 8, 8, 128)	512
conv2d_209 (Conv2D)	(None, 8, 8, 128)	147584
activation_259 (Activation)	(None, 8, 8, 128)	0
batch_normalization_225 (Bat	(None, 8, 8, 128)	512
max_pooling2d_103 (MaxPoolin	(None, 4, 4, 128)	0
dropout_124 (Dropout)	(None, 4, 4, 128)	0
conv2d_210 (Conv2D)	(None, 4, 4, 256)	295168
activation_260 (Activation)	(None, 4, 4, 256)	0
batch_normalization_226 (Bat	(None, 4, 4, 256)	1024

conv2d_211 (Conv2D)	(None, 4, 4, 256)	590080

activation_261 (Activation)	(None, 4, 4, 256)	0

batch_normalization_227 (Batch Normalization)	(None, 4, 4, 256)	1024

max_pooling2d_104 (MaxPooling2D)	(None, 2, 2, 256)	0

dropout_125 (Dropout)	(None, 2, 2, 256)	0

conv2d_212 (Conv2D)	(None, 2, 2, 512)	1180160

activation_262 (Activation)	(None, 2, 2, 512)	0

batch_normalization_228 (Batch Normalization)	(None, 2, 2, 512)	2048

conv2d_213 (Conv2D)	(None, 2, 2, 512)	2359808

activation_263 (Activation)	(None, 2, 2, 512)	0

batch_normalization_229 (Batch Normalization)	(None, 2, 2, 512)	2048

max_pooling2d_105 (MaxPooling2D)	(None, 1, 1, 512)	0

dropout_126 (Dropout)	(None, 1, 1, 512)	0

flatten_25 (Flatten)	(None, 512)	0

dense_50 (Dense)	(None, 128)	65664

activation_264 (Activation)	(None, 128)	0

batch_normalization_230 (Batch Normalization)	(None, 128)	512

dropout_127 (Dropout)	(None, 128)	0

dense_51 (Dense)	(None, 10)	1290

activation_265 (Activation)	(None, 10)	0
=====		
Total params: 4,787,626		
Trainable params: 4,783,402		
Non-trainable params: 4,224		

None		

By testing our model, we ended up having an overall accuracy, recall, and f1-score of 0.84. Generally, most of the error came from not being able to properly distinguish cats and dogs as well as smaller errors coming from having difficulty distinguishing animals in general. The model did relatively better in accurately classifying inanimate objects (ships, trucks, airplanes, and automobiles). The accuracy of predictions could be potentially improved by training the model without these inanimate objects and including more training data of the similar animals (cats and dogs). This might allow the model to find smaller differences between animals that may be of similar sizes or have similar features (like cats and dogs).